

# How do you construct a Knowledge Graph?

Ana Iglesias-Molina and David Chaves-Fraga  
Half-day tutorial at ECAI2024

# About us...

---



**David Chaves-Fraga**  
Assistant Professor  
*CITIUS / GSI*  
Universidade de Santiago de Compostela



**Ana Iglesias-Molina**  
PhD Student  
*Ontology Engineering Group (OEG)*  
Universidad Politécnica de Madrid



# KGC-CG in W3C

---

---



W3C Community Group Website

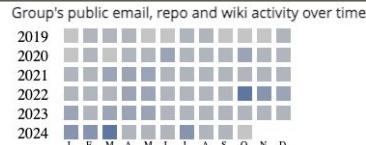


W3C Community Group GitHub

## KNOWLEDGE GRAPH CONSTRUCTION COMMUNITY GROUP

The overall goal of this community group is to support its participants into developing better methods for Knowledge Graphs construction. The Community Group will (i) study current Knowledge Graph construction methods and implementations, (ii) identify the corresponding requirements and issues that hinder broader Knowledge Graph construction, (iii) discuss use cases, (iv) formulate guidelines, best practices and test cases for Knowledge Graph construction, (v) develop methods, resources and tools for evaluating Knowledge Graphs construction, and in general (vi) continue the development of the W3C-recommended R2RML language beyond relational databases. The proposed Community Group could be instrumental to advance research, increase the level of education and awareness and enable learning and participation with respect to Knowledge Graph construction.

[kg-construct](#)



*Note: Community Groups are proposed and run by the community. Although W3C hosts these conversations, the groups do not necessarily represent the views of the W3C Membership or staff.*

Chairs, when logged in, may publish draft and final reports. Please see [report requirements](#).

### biweekly meetings

Anastasia Dimou | Posted on: April 6, 2021

The Community Group meets every 2nd Monday at 15h CET.

The details to join the meetings:

Join Zoom Meeting  
<https://upm.zoom.us/j/87908976407>

### Tools for this group i

Mailing List

IRC

Github repositories

RSS

Contact This Group

### Get involved i

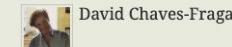
Anyone may join this Community Group. All participants in this group have signed the W3C Community Contributor License Agreement.

#### JOIN OR LEAVE THIS GROUP

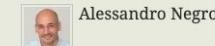


Anastasia Dimou

Chairs  
↗



David Chaves-Fraga



Alessandro Negro

### Participants (184)



[View all participants](#)

# Technical requirements

---

Text Editor with YAML support (e.g., Sublime Text)

Google Colab / Jupyter Notebook

<https://w3id.org/kg-construct/tutorials/ecai2024>

# Agenda

---

- RDF Knowledge Graphs and the Semantic Web
- RML: The RDF Mapping Language
- Coffee Break
- Hands-on session on KG Construction:
  - Mapping Data with YARRRML
  - Transforming Heterogeneous Data to RDF
- Resources for KG Construction
  - Ontology development, mapping supporters, KGC engines, benchmarks

# Knowledge Graphs

14:00 - 14:30



Knowledge Graph definition

Semantic Web technologies

KG Construction

History of mapping languages



# My LinkedIn these days...

---

---



Will Powell  
Design, Innovation and Generative AI for enterprise at Joule

Researchers created a knowledge graph based on 58 million journal papers to fuel personalized research ideas for scientists. Over 4400 ideas generated by the system SciMuse were evaluated by 100 research group leaders.

They were then able to predict with strong accuracy which of its ideas would be ranked most interesting, using a couple of different approaches.

The greatest future potential for science lies in interdisciplinary collaboration. Tools such as SciMuse provide massive potential for novel and unexpected research ideas that might not otherwise come to the surface.

Interestingly, ideas generated that were more central to the knowledge graph, and thus highly connected, were rated as less interesting. However the most interesting ideas tended to be from relatively adjacent rather distant domains.

Knowledge graphs offer an exceptionally powerful approach to uncovering high-potential research directions. It will be exciting to see how their rapidly growing use will shape scientific progress.



Will Powell - Data...

Design, Innovation and Generative AI for enterprise at Joule

Exciting times in the enterprise AI world! SAP announced their Knowledge Graph solution! I created an AI-focused podcast episode explaining what a knowledge graph is. Using NotebookLM and Adobe Express AI Animator.

SAP's Knowledge Graph is set to boost how AI can leverage data and how Joule can understand the SAP data:

- ⌚ Grounds AI in specific business semantics
- 🕒 Connects complex relationships between business entities
- 📊 Reduces manual data modeling complexity
- 🔍 Improves accuracy and relevance of AI outputs

The ability to automatically understand relationships between business objects (like connecting purchase orders, invoices, and customers) marks a significant advancement in building intelligent applications!💡

What are your thoughts on SAP's knowledge graphs? How do you see this impacting the enterprise AI landscape?

#EnterpriseAI #SAP #KnowledgeGraphs #AI #DataInnovation  
#DigitalTransformation

# Who has a KG?

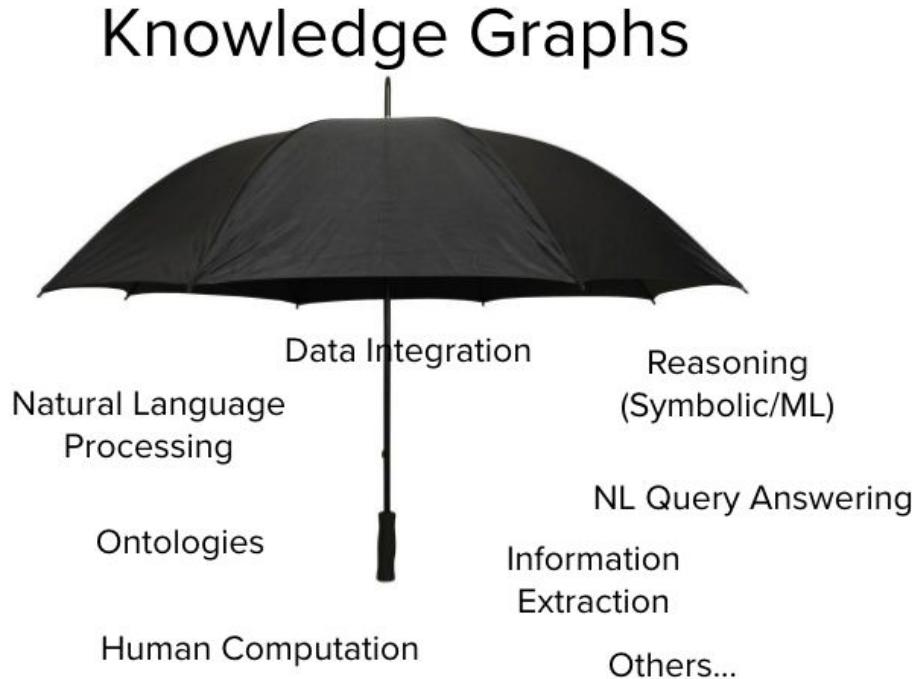
---

---



# What is a Knowledge Graph?

---



# What is a Knowledge Graph?

---

“we define a knowledge graph as a **graph of data intended to accumulate and convey knowledge of the real world**, whose **nodes represent entities** of interest and whose **edges represent potentially different relations between these entities**”



Hogan, A., Blomqvist, E., Cochez, M., d'Amato, C., Melo, G. D., Gutierrez, C., ... & Zimmermann, A. (2021). Knowledge graphs. *ACM Computing Surveys (Csur)*, 54(4), 1-37.

# What is a Knowledge Graph?

---

**Knowledge:** something that is known and can be written down.

Knowledge may be **composed of simple or quantified statements**

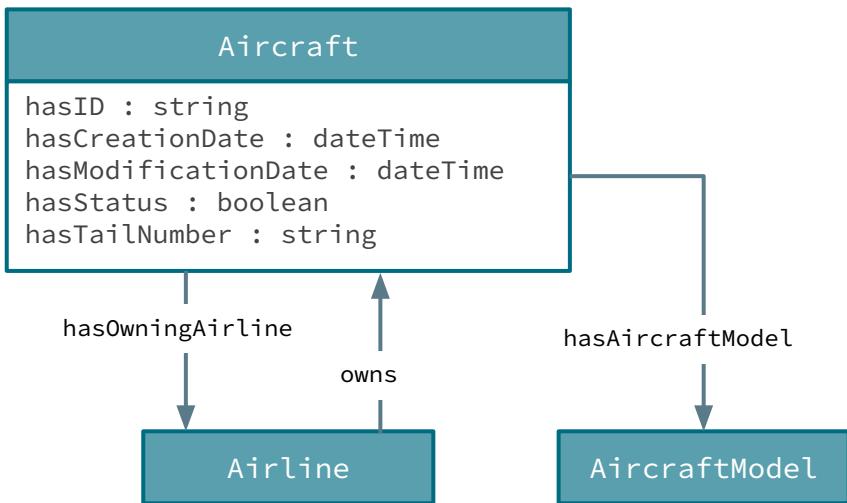
For quantified statements, a more expressive way to represent knowledge –**such as ontologies** or rules– is required



Hogan, A., Blomqvist, E., Cochez, M., d'Amato, C., Melo, G. D., Gutierrez, C., ... & Zimmermann, A. (2021). Knowledge graphs. *ACM Computing Surveys (Csur)*, 54(4), 1-37.

# TBox & ABox

Tbox – Ontology



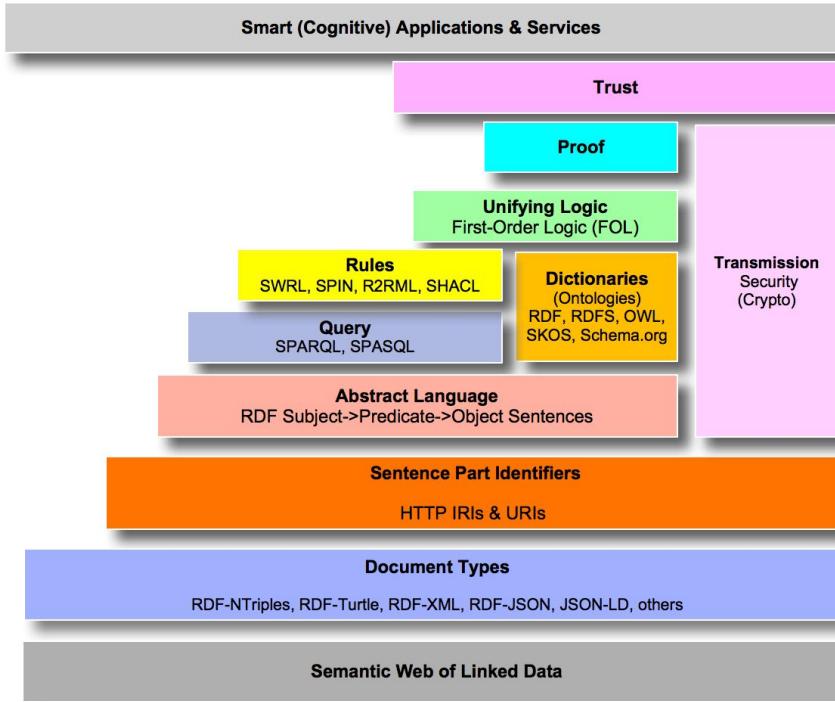
ABox – Data Graph

IB\_X123 is a Aircraft  
IB\_X123 hasID "IBX123"  
IB\_X123 hasCreationDate 28/12/2021  
...  
...  
IB\_X123 hasAircraftModel X123  
X123 is a AircraftModel  
Iberia owns IB\_X123  
Iberia is a Airline

ABox statements must be TBox-compliant:  
they are assertions that use the vocabulary defined by the TBox.

# The Semantic Web stack

---



Ontologies → OWL

Data → RDF

Queries → SPARQL

Validation → SHACL

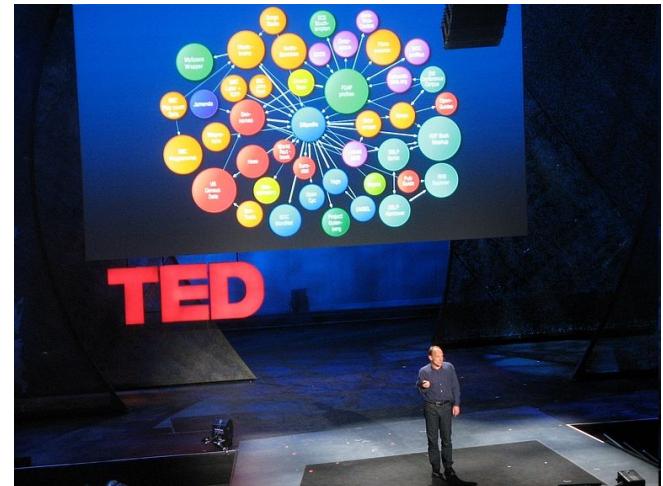
Construction → [R2]RML

# Semantic Web: RDF, OWL, etc.

---

<http://www.w3.org/DesignIssues/LinkedData.html>

- Use URIs as names for things
- Use HTTP URIs so that people can look up those names.
- When someone looks up a URI, provide useful information using standards (RDF\*, SPARQL)
- Include links to other URIs, so that they can discover more things.

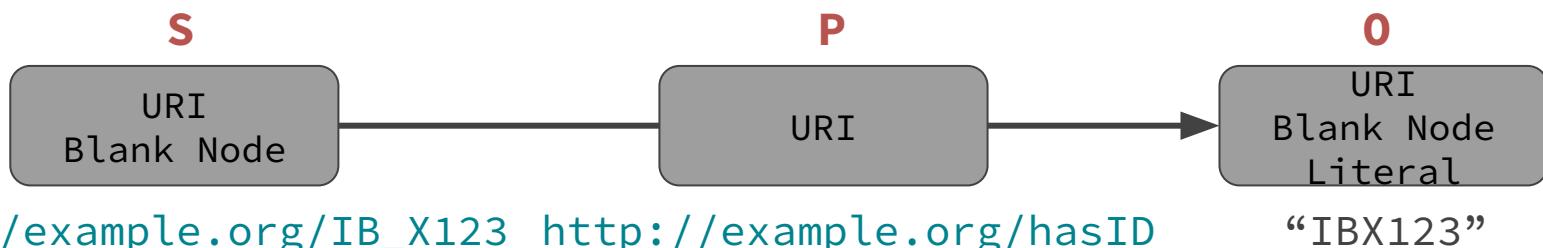


[https://youtu.be/OM6XIIcm\\_qo](https://youtu.be/OM6XIIcm_qo)

# Semantic Web: RDF



RDF triple is composed by (Subject Predicate Object), where:



RDF graph is a **set** of RDF **triples**.

An RDF graph can be identified by an URI (RDF Quads)

A set of RDF graphs is an RDF Dataset

**Literals** may have an XSD datatype associated

# Semantic Web: RDF

---

IB\_X123 is a Aircraft  
IB\_X123 hasID "IBX123"  
IB\_X123 hasCreationDate 28/12/2021  
IB\_X123 hasAircraftModel X123  
X123 is a AircraftModel  
Iberia owns IB\_X123  
Iberia is a Airline

```
@prefix ex: <http://example-ecai.org/resource/>
@prefix onto: <http://example-ecai.org/ontology/>
@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>
@prefix xsd: <http://www.w3.org/2001/XMLSchema#>
```

ex:IB_X123	rdf:type	onto:Aircraft
ex:IB_X123	onto:hasID	"IBX123"^^xsd:string
ex:IB_X123	onto:hasCreationDate	"2021-12-28"^^xsd:date
ex:IB_X123	onto:hasAircraftModel	ex:X123
ex:X123	rdf:type	onto:AircraftModel
ex:Iberia	onto:owns	ex:IB_X123
ex:Iberia	rdf:type	onto:Airline





# Semantic Web: Ontologies - OWL

OWL / RDFS

owl:ObjectProperty      owl:DataProperty

rdf:type

rdf:type

owl:Class

rdf:subClassOf

Vocabulary

onto:Aircraft

onto:hasID : string  
onto:hasCreationDate : dateTime  
onto:hasModificationDate : dateTime  
onto:hasStatus : boolean  
onto:hasTailNumber : string

onto:Airline

onto:hasAirline

onto:owns

onto:AircraftModel

onto:TransportType

Data

ex:Iberia

rdf:type

ex:IB\_X123

onto:owns

rdf:type

ex:X123

onto:hasAirline

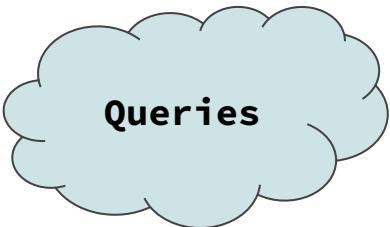
onto:hasModel

rdf:type

# Semantic Web: SPARQL & SHACL

---

---



## SPARQL 1.1 Overview

W3C Recommendation 21 March 2013

**This version:**

<https://www.w3.org/TR/2013/REC-sparql11-overview-20130321/>

**Latest version:**

<https://www.w3.org/TR/sparql11-overview/>

**Previous version:**

<https://www.w3.org/TR/2012/PR-sparql11-overview-20121108/>

**Editor:**

The W3C SPARQL Working Group, see [Acknowledgements <public-rdf-dawg-comments@w3.org](mailto:Acknowledgements<public-rdf-dawg-comments@w3.org)

Please refer to the [errata](#) for this document, which may include some normative corrections.

See also [translations](#).

Copyright © 2013 W3C® (MIT, ERCIM, Keio, Beihang), All Rights Reserved. W3C liability, trademark and [document use](#) rule

## Shapes Constraint Language (SHACL)

W3C Recommendation 20 July 2017



**This version:**

<https://www.w3.org/TR/2017/REC-shacl-20170720/>

**Latest published version:**

<https://www.w3.org/TR/shacl/>

**Latest editor's draft:**

<https://w3c.github.io/data-shapes/shacl/>

**Implementation report:**

<https://w3c.github.io/data-shapes/data-shapes-test-suite/>

**Previous version:**

<https://www.w3.org/TR/2017/PR-shacl-20170608/>

**Editors:**

Holger Knublauch, TopQuadrant, Inc.

Dimitris Kontokostas, University of Leipzig

**Repository:**

[GitHub](#)

[Issues](#)

**Test Suite:**

[SHACL Test Suite](#)

Please check the [errata](#) for any errors or issues reported since publication.

See also [translations](#).





# Semantic Web: Construction with R2RML



## R2RML: RDB to RDF Mapping Language

W3C Recommendation 27 September 2012

This version:

<http://www.w3.org/TR/2012/REC-r2rml-20120927/>

Latest version:

<http://www.w3.org/TR/r2rml/>

Previous version:

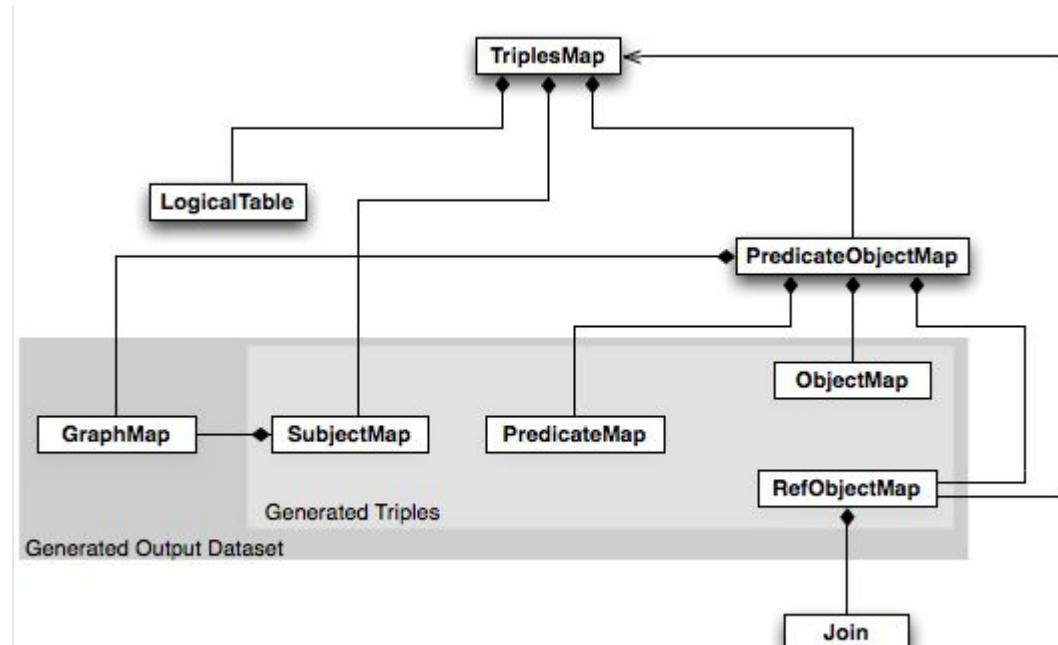
<http://www.w3.org/TR/2012/PR-r2rml-20120814/>

Editors:

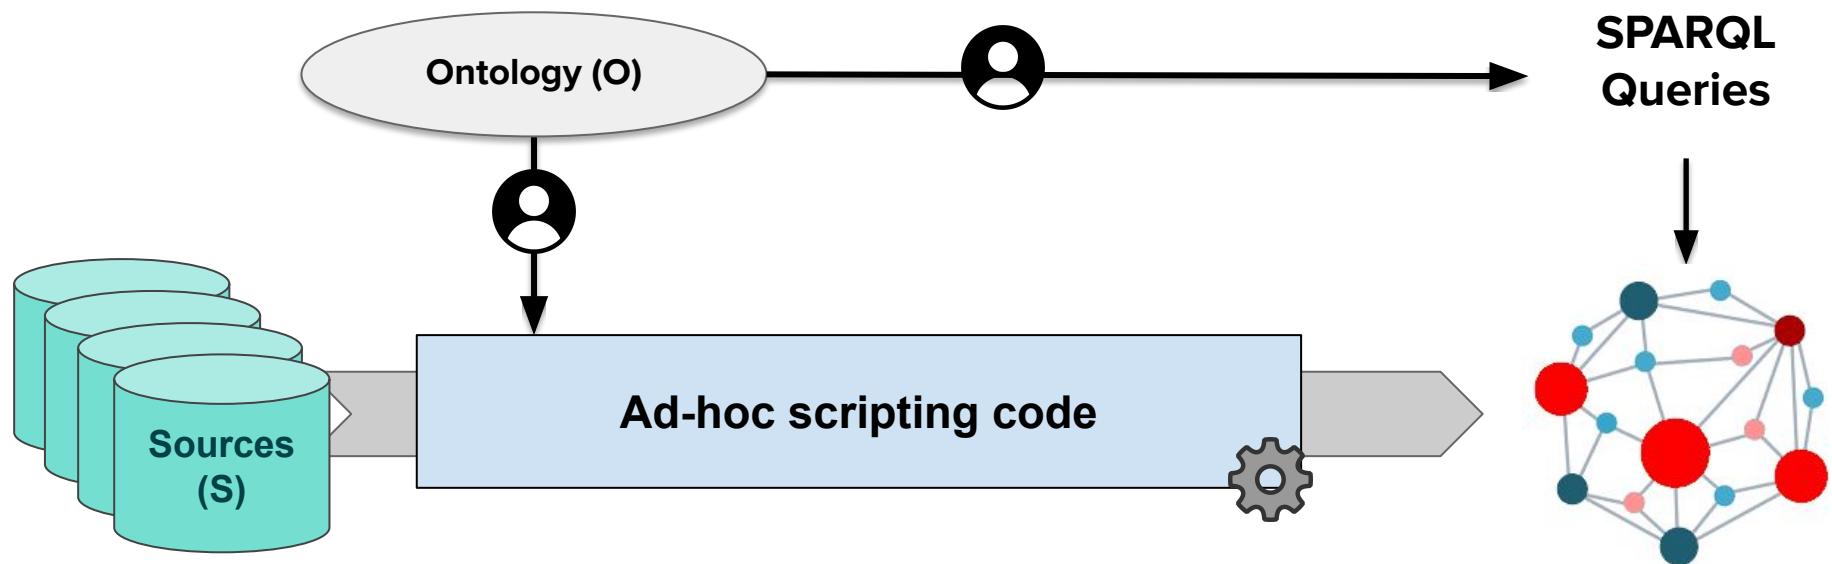
Souripriya Das, Oracle  
Seema Sundara, Oracle  
Richard Cyganiak, DERI, National University of Ireland, Galway

Please refer to the [errata](#) for this document, which may include some normative

See also [translations](#).



# Knowledge Graph Construction: Script based



Efficient



Scalable



Maintainable



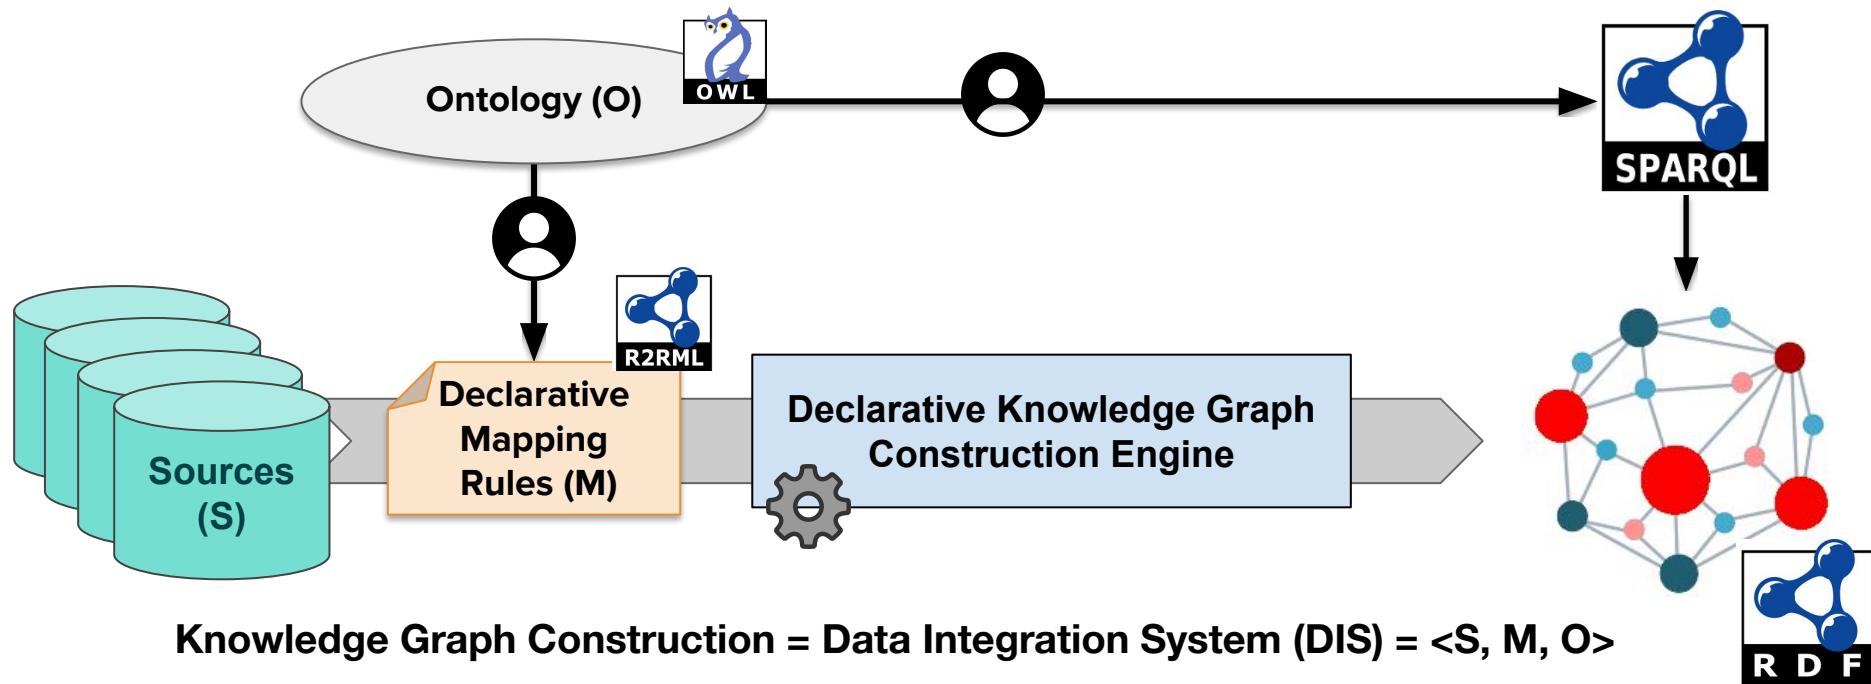
Robust



Reproducible

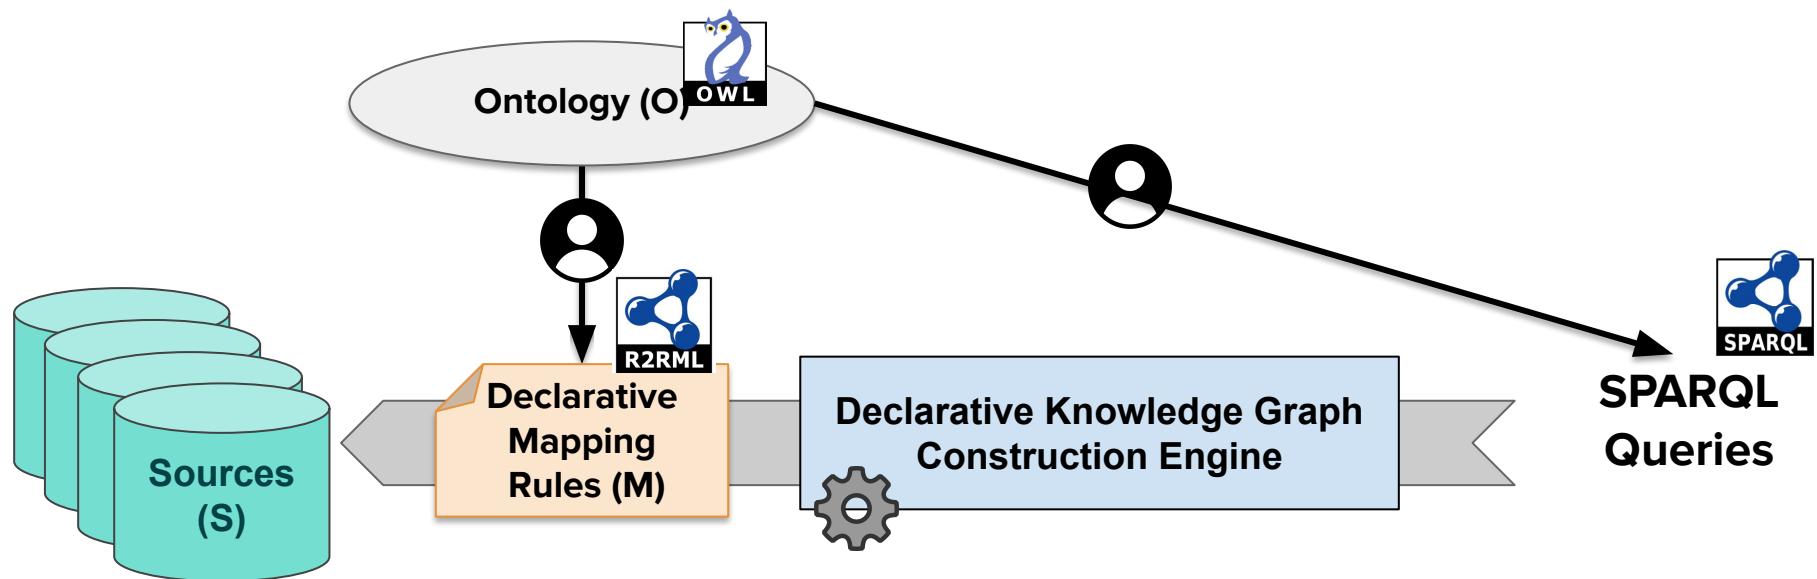


# Knowledge Graph Construction: Materialization



Poggi, A., Lembo, D., Calvanese, D., De Giacomo, G., Lenzerini, M., & Rosati, R. (2008). Linking data to ontologies. In *Journal on data semantics X*  
Lenzerini, M. Data integration: A theoretical perspective. In *Proceedings of the 21st ACM SIGMOD-SIGACT-SIGART symposium on Principles of database systems*

# Knowledge Graph Construction: Virtualization



**Knowledge Graph Construction = Data Integration System (DIS) =  $\langle S, M, O \rangle$**



Poggi, A., Lembo, D., Calvanese, D., De Giacomo, G., Lenzerini, M., & Rosati, R. (2008). Linking data to ontologies. In *Journal on data semantics X*  
Lenzerini, M. Data integration: A theoretical perspective. In *Proceedings of the 21st ACM SIGMOD-SIGACT-SIGART symposium on Principles of database systems*

# Mapping Languages for KG Construction

14:30 - 15:30

Mapping languages and evolution

RML: The RDF Mapping Language



— — —

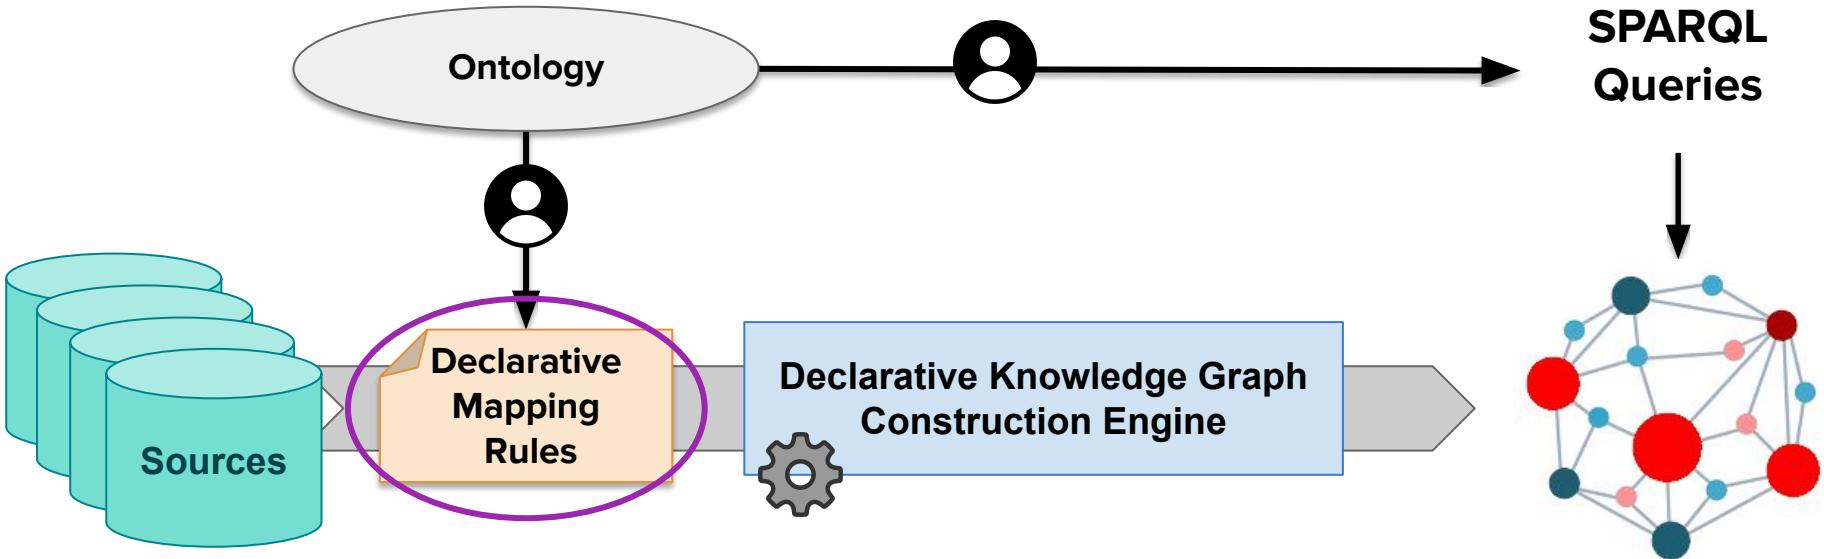
# Who is using RML

---

- Orange
- EU Comission
- Bosch
- Google
- BASF
-

# Mapping rules

---



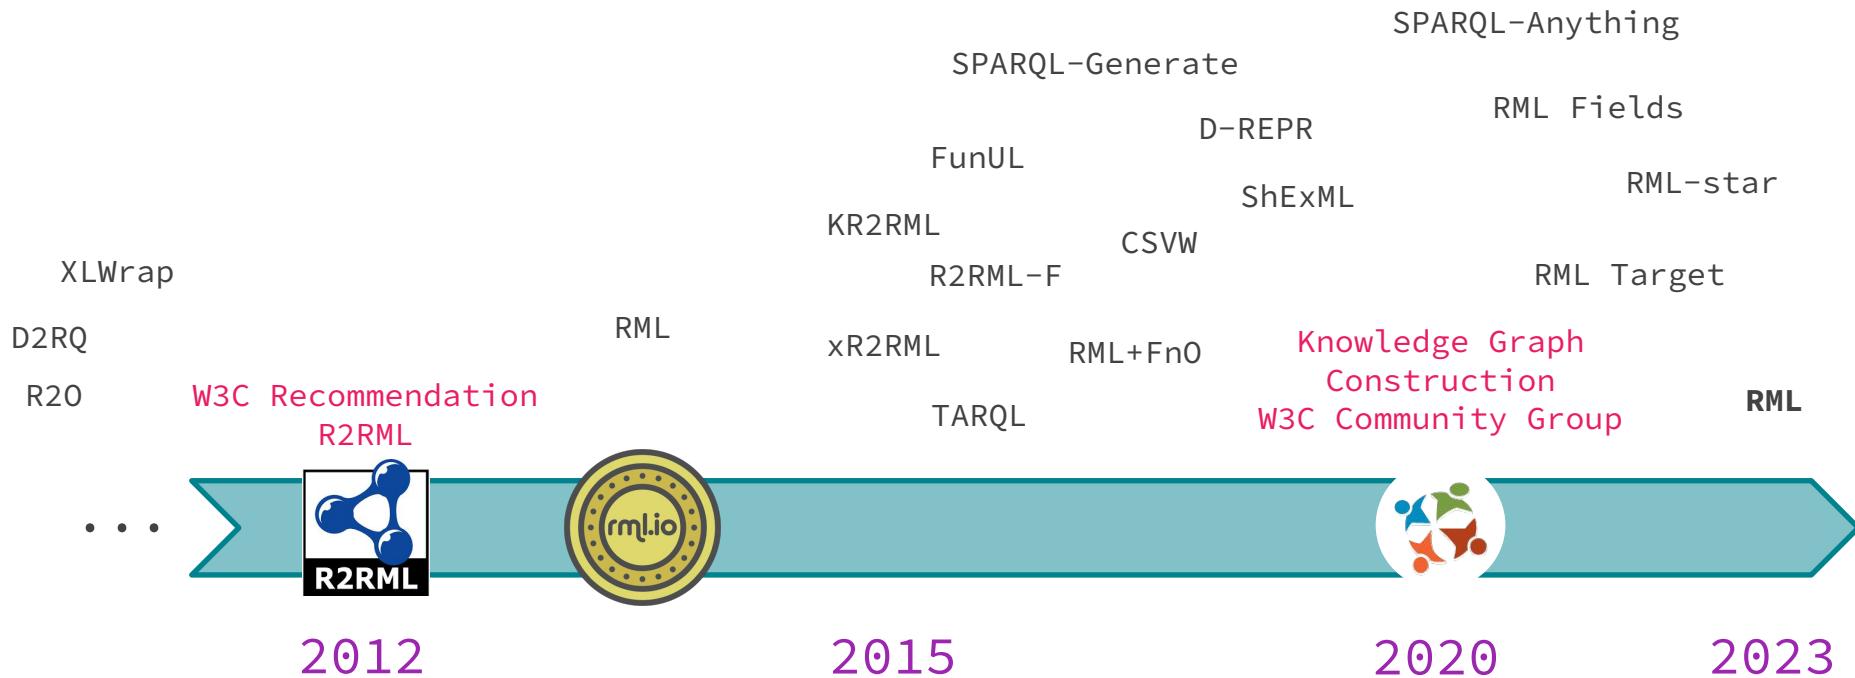
# Mapping languages

---

- Allow specifying **transformation rules** from heterogeneous data sources into RDF graphs
  - **Data source** and **access** description
  - **Triple** generation:
    - Subject, Predicate, Object
    - Language tags and datatypes
  - Additional:
    - Join conditions
    - Data transformation functions
- Many different languages proposed throughout the years
  - Different **features**
  - Variety of **implementations**

# The history so far

---



# Mapping example

—

## RML

## Data sources

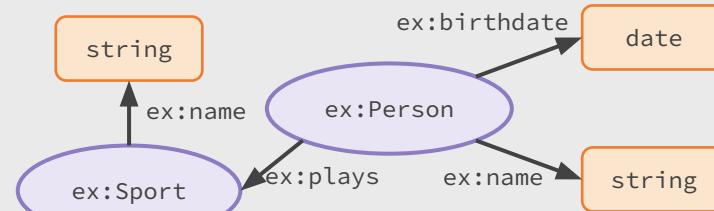
people.csv

ID	Name	Age	SportID
1	Emily	23	2
2	Jonah	22	2

sports.csv

ID	Sport
1	Ice Skating
2	Rugby

## Ontology



## Mapping



<#PersonTM>

Group of  
transformation rules

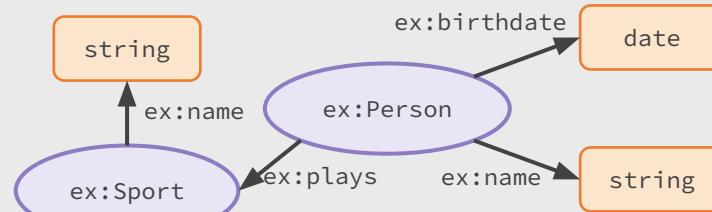
## Output Knowledge Graph

## Data sources

people.csv			
ID	Name	Age	SportID
1	Emily	23	2
2	Jonah	22	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



## Mapping



```
<#PersonTM>  
rml:logicalSource [  
    rml:source "people.csv" ;  
    rml:referenceFormulation ql:CSV ];
```

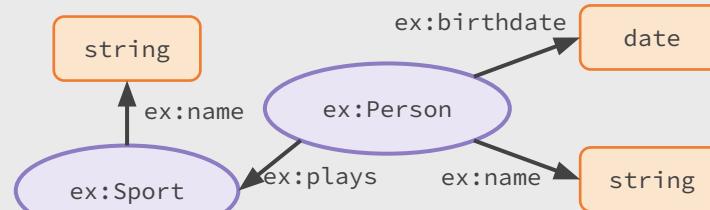
## Output Knowledge Graph

## Data sources

people.csv			
ID	Name	Age	SportID
1	Emily	23	2
2	Jonah	22	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



## Mapping



```

<#PersonTM>
rml:logicalSource [
  rml:source "people.csv" ;
  rml:referenceFormulation ql:CSV ];
rml:subjectMap [
  rml:template "Person/{ID}";
  rml:class ex:Person ];
  
```

## Output Knowledge Graph

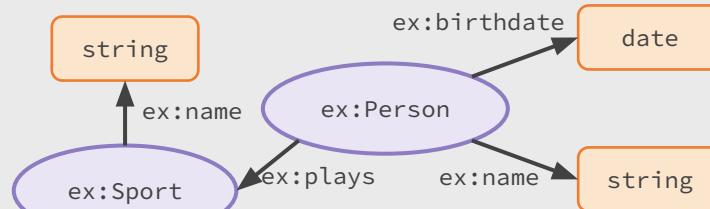


## Data sources

people.csv			
ID	Name	Age	SportID
1	Emily	23	2
2	Jonah	22	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



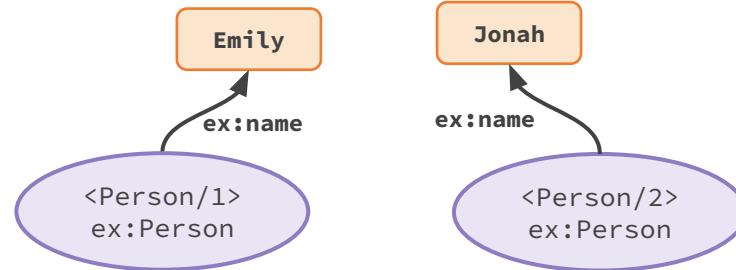
## Mapping



```

<#PersonTM>
rml:logicalSource [
  rml:source "people.csv" ;
  rml:referenceFormulation ql:CSV ];
rml:subjectMap [
  rml:template "Person/{ID}";
  rml:class ex:Person ];
rml:predicationObjectMap [
  rml:predication ex:name;
  rml:objectMap [
    rml:reference "Name" ]];
  
```

## Output Knowledge Graph

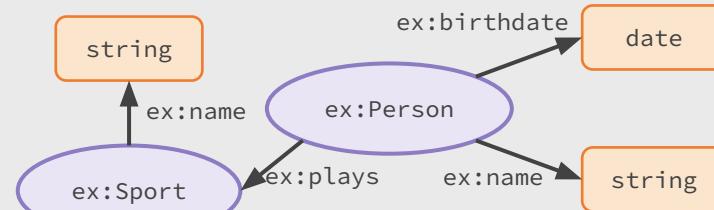


## Data sources

people.csv			
ID	Name	Age	SportID
1	Emily	23	2
2	Jonah	22	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



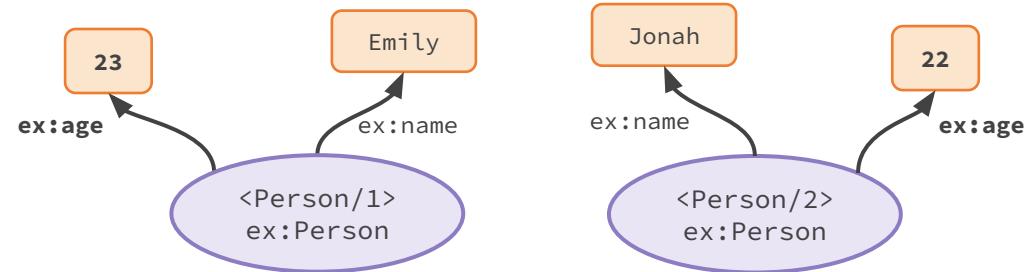
## Mapping



```

<#PersonTM>
rml:logicalSource [
  rml:source "people.csv" ;
  rml:referenceFormulation ql:CSV ];
rml:subjectMap [
  rml:template "Person/{ID}";
  rml:class ex:Person ];
rml:predicateObjectMap [
  rml:predicate ex:name;
  rml:objectMap [
    rml:reference "Name" ];
rml:predicateObjectMap [
  rml:predicate ex:age;
  rml:objectMap [
    rml:reference "Age" ];
...
  
```

## Output Knowledge Graph

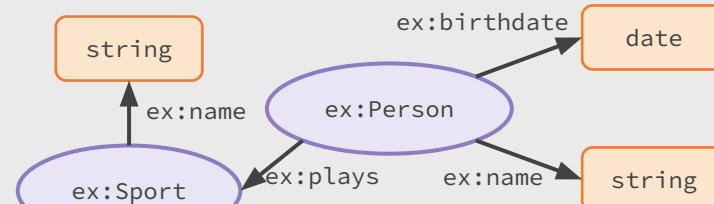


## Data sources

people.csv			
ID	Name	Age	SportID
1	Emily	23	2
2	Jonah	22	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



## Mapping



```

<#PersonTM>
rml:logicalSource [
  rml:source "people.csv" ;
  rml:referenceFormulation ql:CSV ];
rml:subjectMap [
  rml:template "Person/{ID}";
  rml:class ex:Person ];
rml:predicateObjectMap [
  rml:predicate ex:name;
  rml:objectMap [
    rml:reference "Name" ];
  rml:predicateObjectMap [
    rml:predicate ex:birthdate;
    rml:objectMap [
      rml:reference "Birthdate"]
  ...
]
  
```

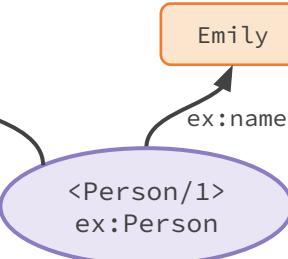
### <#SportTM>

```

rml:logicalSource [
  rml:source "sports.csv" ;
  rml:referenceFormulation ql:CSV ];
rml:subjectMap [
  rml:template "Sport/{ID}";
  rml:class ex:Sport ];
rml:predicateObjectMap [
  rml:predicate ex:name;
  rml:objectMap [
    rml:reference "Sport"]].
  
```

23

ex:age



Emily

ex:name

Jonah

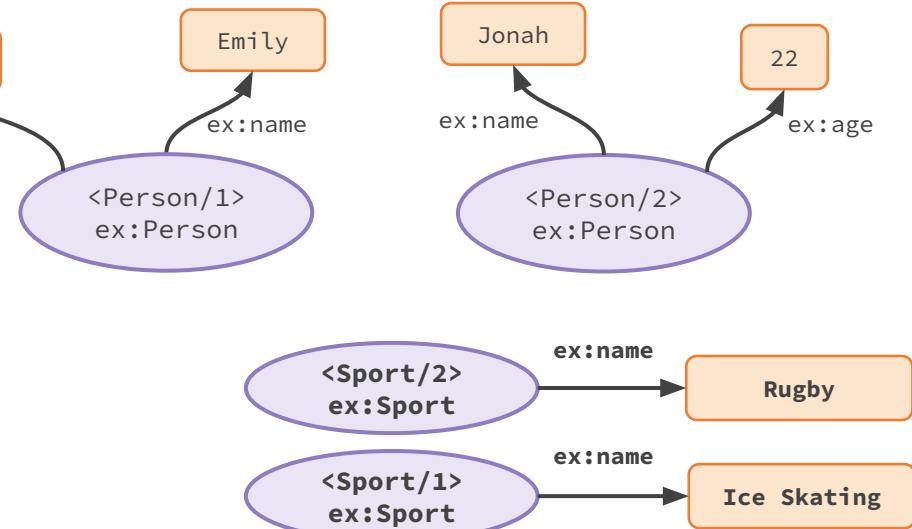
ex:name

22

ex:age



## Output Knowledge Graph

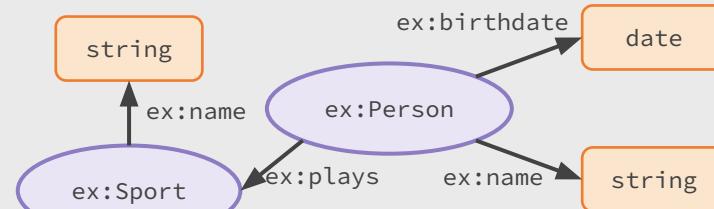


## Data sources

people.csv			
ID	Name	Age	SportID
1	Emily	23	2
2	Jonah	22	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



## Mapping



```

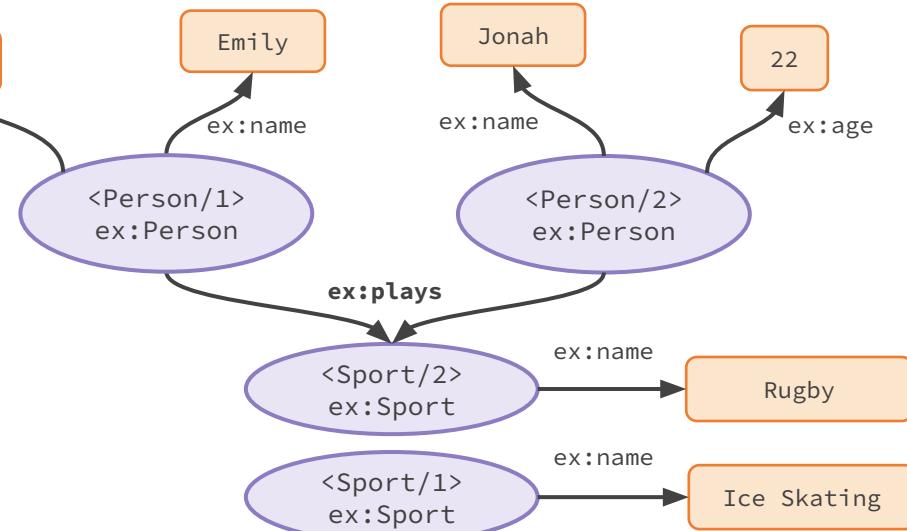
<#PersonTM>
...
rml:predicatesObjectMap [
  rml:predicates ex:plays;
  rml:objectMaps [
    rml:parentTriplesMap
      <#SportTM>;
    rml:joinCondition [
      rml:child "SportID";
      rml:parent "ID";
    ];
  ];
];
  
```

### <#SportTM>

```

rml:logicalSource [
  rml:source "sports.csv";
  rml:referenceFormulation ql:CSV];
rml:subjectMap [
  rml:template "Sport/{ID}";
  rml:class ex:Sport];
rml:predicatesObjectMap [
  rml:predicates ex:name;
  rml:objectMaps [
    rml:reference "Sport"]].
  
```

## Output Knowledge Graph

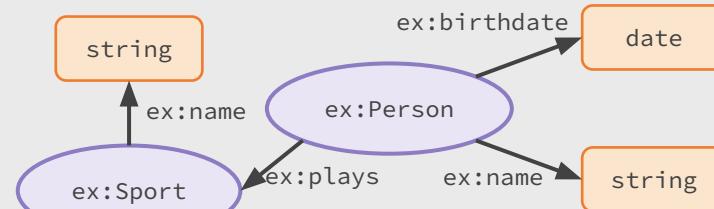


## Data sources

people.csv			
ID	Name	Age	SportID
1	Emily	23	2
2	Jonah	22	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



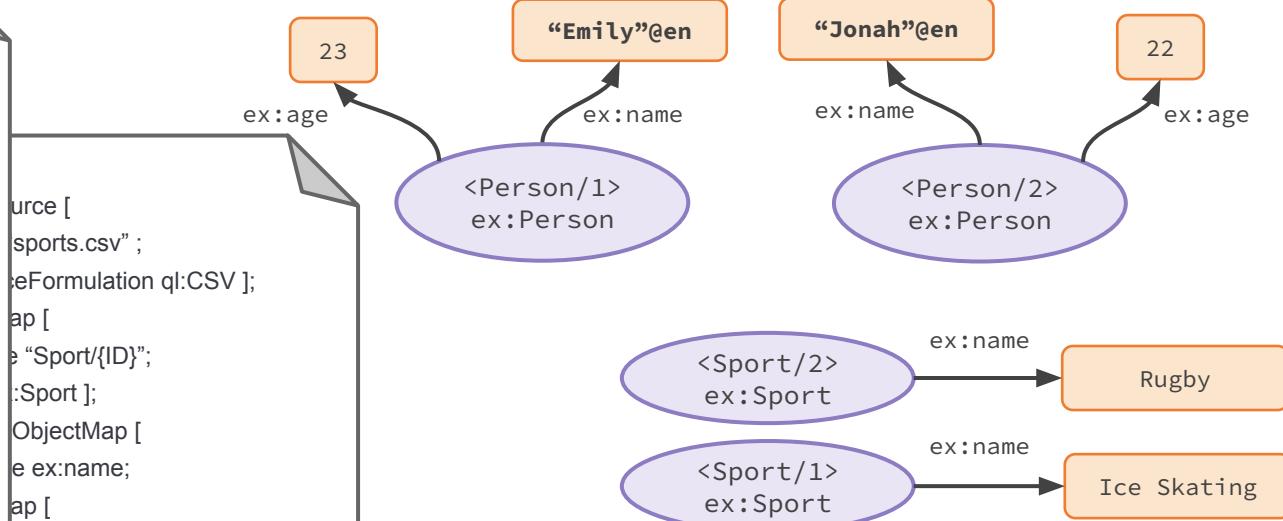
## Mapping



```

<#PersonTM>
rml:logicalSource [
  rml:source "people.csv" ;
  rml:referenceFormulation ql:CSV ];
rml:subjectMap [
  rml:template "Person/{ID}";
  rml:class ex:Person ];
rml:predicateObjectMap [
  rml:predicate ex:name;
  rml:objectMap [
    rml:reference "Name"; rml:language "en" ]];
rml:predicateObjectMap [
  rml:predicate ex:birthdate;
  rml:objectMap [
    rml:reference "Birthdate"]
...
  
```

## Output Knowledge Graph

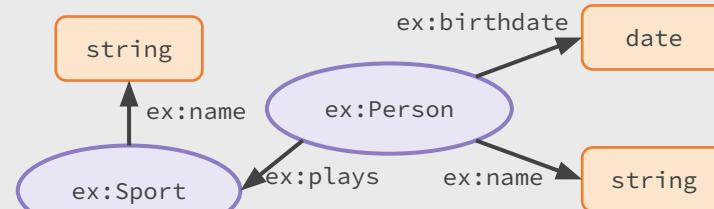


## Data sources

people.csv			
ID	Name	Age	SportID
1	Emily	23	2
2	Jonah	22	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



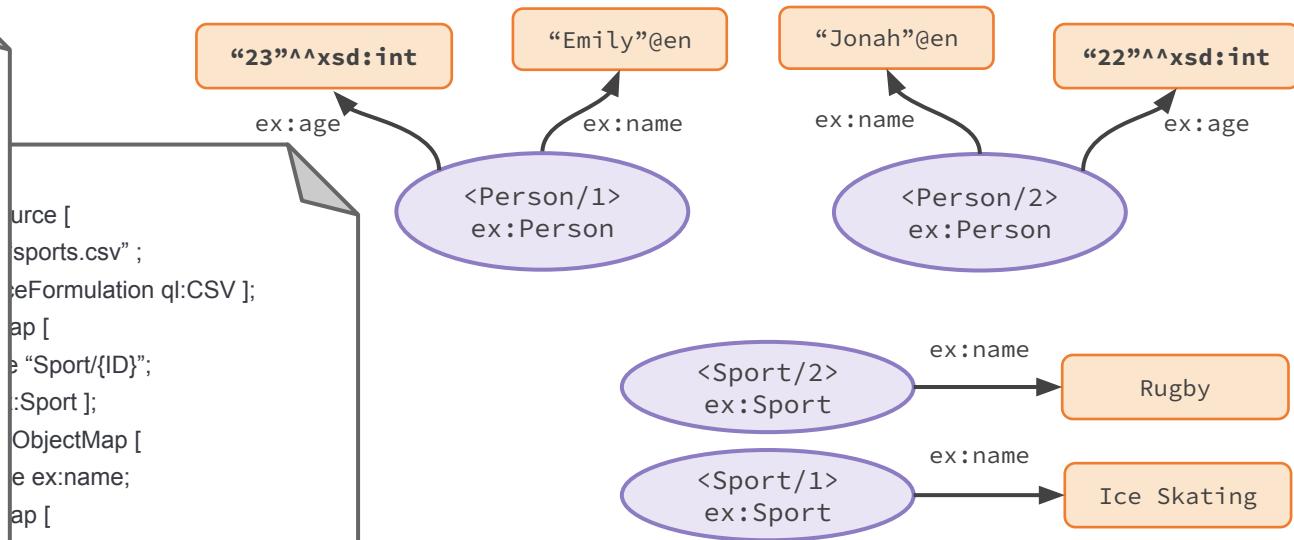
## Mapping



```

<#PersonTM>
rml:logicalSource [
  rml:source "people.csv" ;
  rml:referenceFormulation ql:CSV ];
rml:subjectMap [
  rml:template "Person/{ID}";
  rml:class ex:Person ];
rml:predicateObjectMap [
  rml:predicate ex:name;
  rml:objectMap [
    rml:reference "Name"; rml:language "en" ]];
rml:predicateObjectMap [
  rml:predicate ex:birthdate;
  rml:objectMap [
    rml:reference "Birthdate"; rml:datatype xsd:int ]];
...
  
```

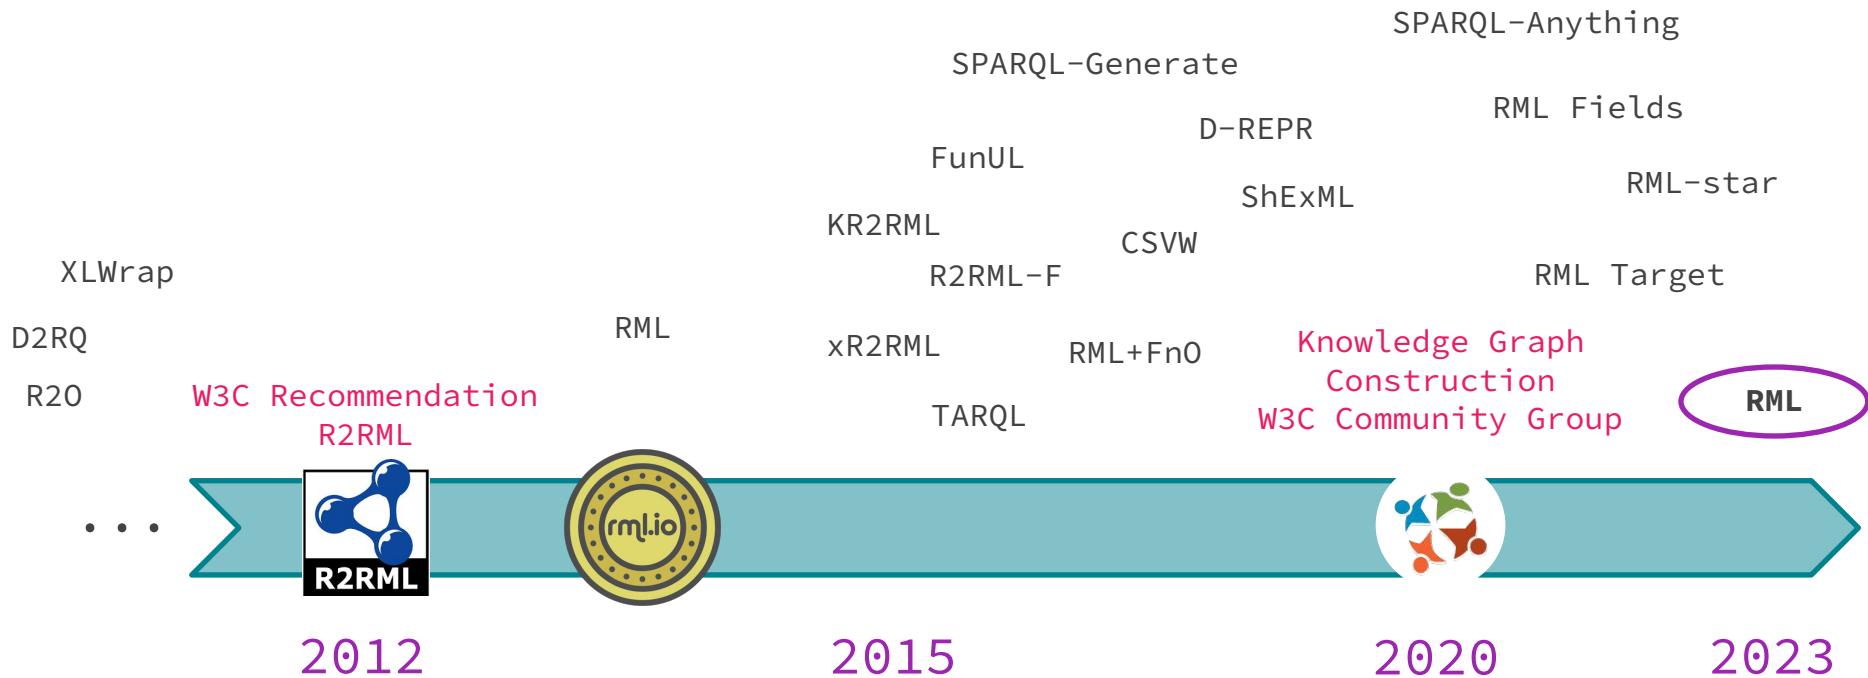
## Output Knowledge Graph



# RML - 2023 revision

# The history so far

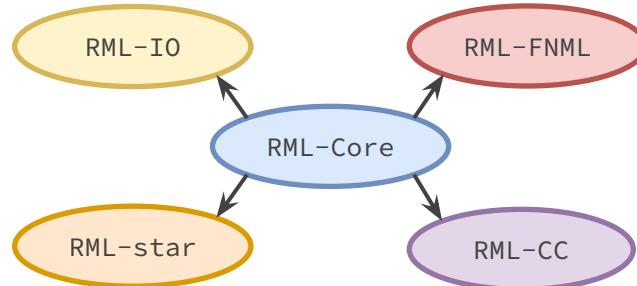
---



# The RML Ontology: Overview

---

- **Reference** ontology specification addressing **identified challenges**
- **Modular** design
- Developed following the **Linked Open Terms (LOT)** methodology
- Maintaining **backwards compatibility** with R2RML
  - **Semantic correspondences** with previous resources
- All resources in all modules use **same namespace and prefix**
  - **rml:** <https://w3id.org/rml/>



# The RML Ontology: Modules

---

---

[Ontologies](#) [Guidelines](#) [Backwards Compatibility](#)

## RML Ontology Modules

Here you can find the list of modules of the mapping language RML.

Ontology	Serialization	License	Language	Links	Description
RML-Core	<a href="#">rdf+xml</a> <a href="#">tar</a>	<a href="#">CC-BY</a>	<a href="#">en</a>	<a href="#">Repository</a> <a href="#">Issues</a> <a href="#">Requirements</a> <a href="#">W3C Specification</a> <a href="#">Shapes</a>	Core ontology that defines the necessary resources to create a mapping.
RML-IO: Source and Target	<a href="#">rdf+xml</a> <a href="#">tar</a>	<a href="#">CC-BY</a>	<a href="#">en</a>	<a href="#">Repository</a> <a href="#">Issues</a> <a href="#">Requirements</a> <a href="#">W3C Specification</a> <a href="#">Shapes</a>	Ontology module that allows the description of input data sources and target outputs.
RML-CC: Collections and Containers	<a href="#">rdf+xml</a> <a href="#">tar</a>	<a href="#">CC-BY</a>	<a href="#">en</a>	<a href="#">Repository</a> <a href="#">Issues</a> <a href="#">Requirements</a> <a href="#">W3C Specification</a> <a href="#">Shapes</a>	Ontology module that allows the generation of collections and containers.
RML-FNML: Functions	<a href="#">rdf+xml</a> <a href="#">tar</a>	<a href="#">CC-BY</a>	<a href="#">en</a>	<a href="#">Repository</a> <a href="#">Issues</a> <a href="#">Requirements</a> <a href="#">W3C Specification</a> <a href="#">Shapes</a>	Ontology module that allows the application of data transformation functions.
RML-Star	<a href="#">rdf+xml</a> <a href="#">tar</a>	<a href="#">CC-BY</a>	<a href="#">en</a>	<a href="#">Repository</a> <a href="#">Issues</a> <a href="#">Requirements</a> <a href="#">W3C Specification</a> <a href="#">Shapes</a>	Ontology module that allows the construction of RDF-star graphs.

Composed of **5 modules**:

- **RML-Core:** schema transformations
- **RML-IO:** source and target data
- **RML-FNML:** data transformation functions
- **RML-CC:** collections and containers
- **RML-star:** RDF-star



[w3id.org/rml/portal](http://w3id.org/rml/portal)

# Network of Resources

---

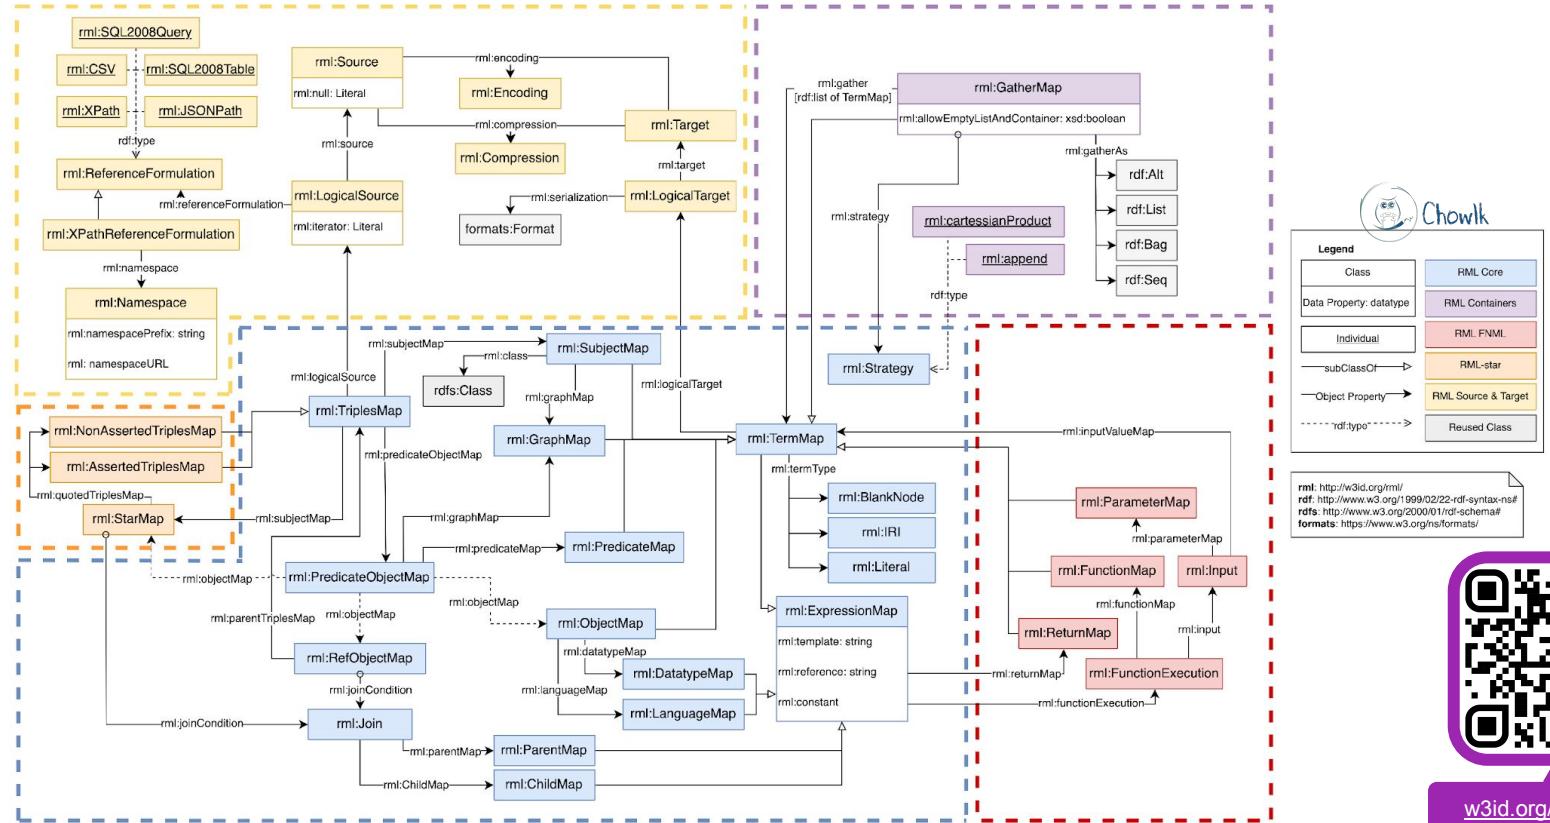
---

Each module is composed of:

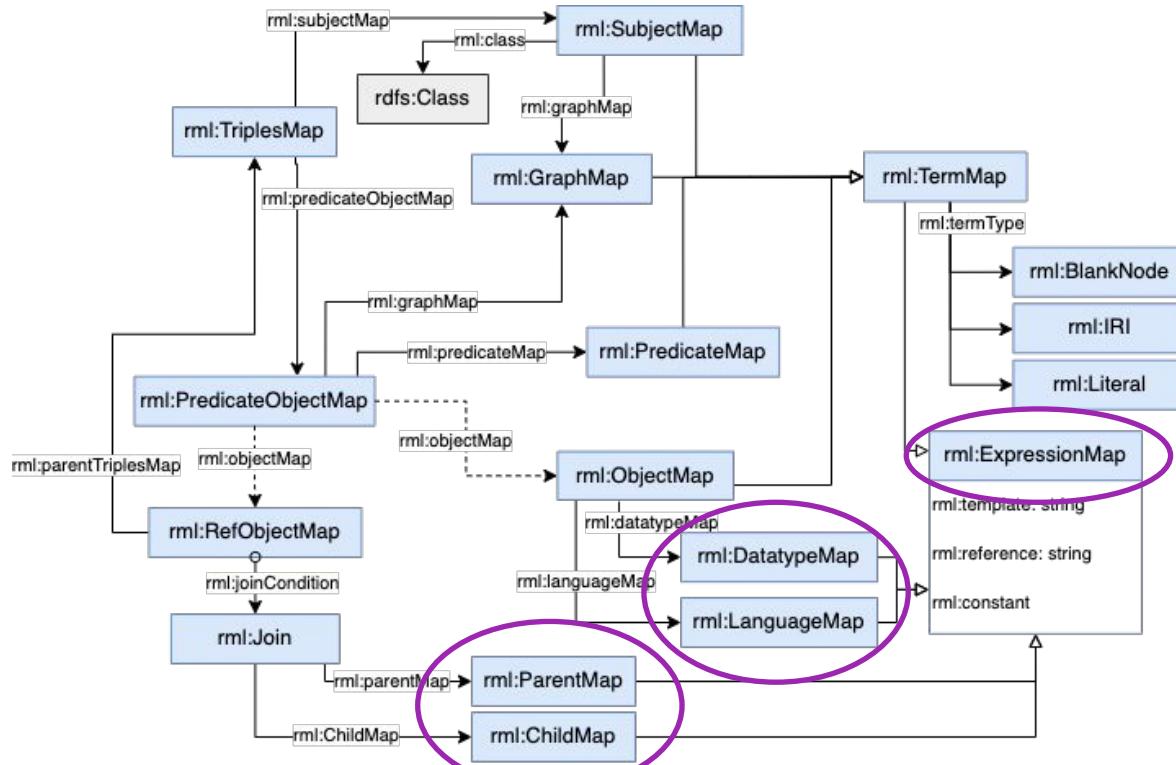
- **Ontology** [w3id.org/rml/core](https://w3id.org/rml/core) [w3id.org/rml/io](https://w3id.org/rml/io) [w3id.org/rml/fnml](https://w3id.org/rml/fnml) [w3id.org/rml/cc](https://w3id.org/rml/cc) [w3id.org/rml/star](https://w3id.org/rml/star)
  - Implemented in OWL2: Definition of concepts, relationships (w/ domain and range)
  - Available in multiple serializations and a human readable documentation
- **Specification** [w3id.org/rml/core/spec](https://w3id.org/rml/core/spec) [w3id.org/rml/io/spec](https://w3id.org/rml/io/spec) [w3id.org/rml/fnml/spec](https://w3id.org/rml/fnml/spec) [w3id.org/rml/cc/spec](https://w3id.org/rml/cc/spec) [w3id.org/rml/star/spec](https://w3id.org/rml/star/spec)
  - Feature details, examples of use and implementation instructions
- **SHACL shapes** [w3id.org/rml/core/shapes](https://w3id.org/rml/core/shapes) [w3id.org/rml/io/shapes](https://w3id.org/rml/io/shapes) [w3id.org/rml/fnml/shapes](https://w3id.org/rml/fnml/shapes) [w3id.org/rml/cc/shapes](https://w3id.org/rml/cc/shapes) [w3id.org/rml/star/shapes](https://w3id.org/rml/star/shapes)
  - Restrictions of module constructs
  - Usable to test correctness of mappings
- **Test cases (WIP)**
  - Language coverage by processing engines



# The RML Ontology: Overview



# RML-Core: Schema transformations



- Maintains **R2RML** basic structure
- **Dynamic** generation of:
  - Language tags
  - Data types
- Increased **flexibility** for join conditions



[w3id.org/rml/core](http://w3id.org/rml/core)

[w3id.org/rml/core/spec](http://w3id.org/rml/core/spec)

[w3id.org/rml/core/shapes](http://w3id.org/rml/core/shapes)

# RML-Core: Example

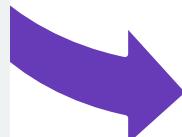
---

```
<#NameTriplesMap> a rml:TriplesMap ;
  rml:logicalSource <#JSONSource> ;
  rml:subjectMap [
    a rml:SubjectMap, rml:ExpressionMap ;
    rml:template "{$.NAME}" ;
    rml:class ex:Athlete ;
  ] ;
  rml:predicateObjectMap [
    a rml:PredicateObjectMap ;
    rml:predicate ex:name ;
    rml:objectMap [
      a rml:ObjectMap, rml:ExpressionMap ;
      rml:reference "$.NAME" ;
      rml:languageMap [
        a rml:LanguageMap;
        rml:reference "$.COUNTRY"]
    ] ] .
```



{JSON}

```
[ { "NAME": "Duplantis",
  "RANK": "1",
  "MARK": "6.22",
  "COUNTRY": "sv"
}, {
  "NAME": "Guttormsen",
  "RANK": "2",
  "MARK": "6.00",
  "COUNTRY": "no" } ]
```

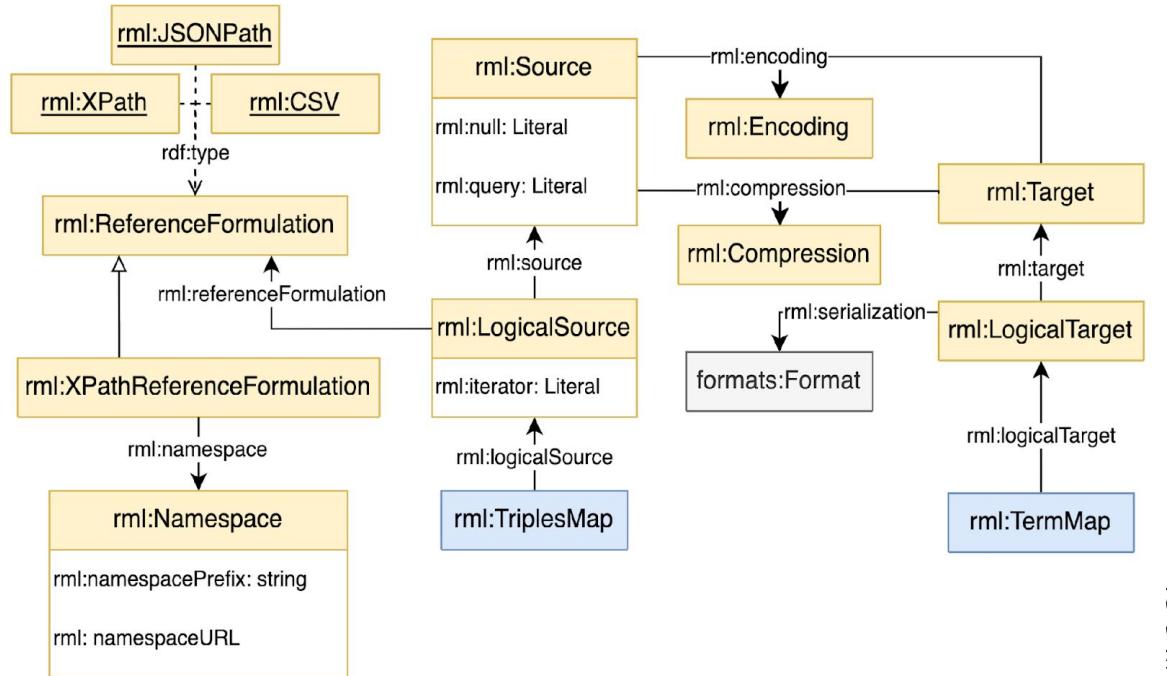


RDF

```
:Duplantis a ex:Athlete ;
  ex:name "Duplantis"@sv .
:Guttormsen a ex:Athlete ;
  ex:name "Guttormsen"@no .
```



# RML-IO: Data source and target



- **Extended input** data source description
- **Output data** description
- Leverage of existing **vocabularies** (SCAT, SPARQL-SD, VoID)



[w3id.org/rml/io](http://w3id.org/rml/io)

[w3id.org/rml/io/spec](http://w3id.org/rml/io/spec)

[w3id.org/rml/io/shapes](http://w3id.org/rml/io/shapes)

# RML-IO: Example

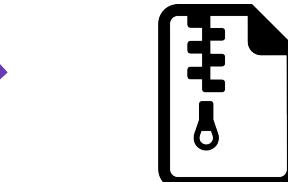
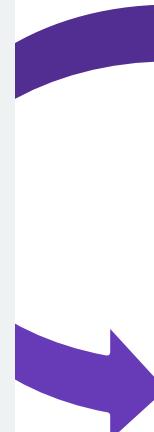
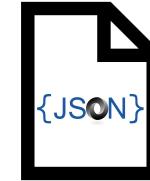
---

```
<#JSONSource> a rml:LogicalSource ;
  rml:source [ a rml:Source, dcat:Distribution ;
    dcat:accessURL <file:///data/ranks.json> ];
  rml:referenceFormulation rml:JSONPath ;
  rml:iterator "$.[*]"; .

<#FileTarget> a rml:LogicalTarget ;
  rml:target [ a rml:Target, void:Dataset ;
    void:dataDump <file:///data/dump.ttl.gz> ;
    rml:compression rml:gzip ;
    rml:encoding rml:UTF-8 ] ;
  rml:serialization formats:Turtle ; .

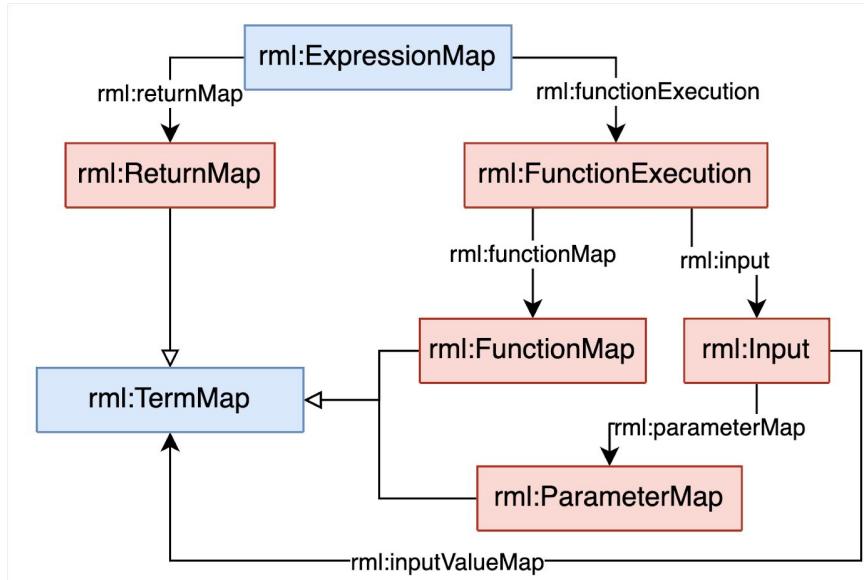
<#RankTriplesMap> a rml:TriplesMap ;
  rml:logicalSource <#JSONSource> ;
  rml:subjectMap <#RankSubjectMap> .

<#RankSubjectMap> rml:logicalTarget <#FileTarget>.
```



# RML-FNML: Data Transformations

---



- Refines **RML+Fn0** approach
- Reference connector between RML and the **Function Ontology** (Fn0)



[w3id.org/rml/fnml](http://w3id.org/rml/fnml)

[w3id.org/rml/fnml/spec](http://w3id.org/rml/fnml/spec)

[w3id.org/rml/fnml/shapes](http://w3id.org/rml/fnml/shapes)

# RML-FNML: Example

```
<#DateTriplesMap> a rml:TriplesMap ;
  rml:logicalSource <#JSONSource> ;
  rml:subjectMap <#RankSubjectMap> ;
  rml:predicateObjectMap [
    rml:predicate ex:jumpOnDate ;
    rml:objectMap [
      rml:functionExecution <#Execution> ;
      rml:return ex:dateOut ] ] .

<#Execution> a rml:FunctionExecution ;
  rml:function ex:parseDate ;
  rml:input
    [ a rml:Input ;
      rml:parameter ex:valueParam;
      rml:inputValueMap [
        rml:reference "$.DATE" ] ] ,
    [ a rml:Input ;
      rml:parameter ex:dateFormatParam ;
      rml:inputValueMap [
        rml:constant "DD-MM-YYYY" ] ] .
```

{JSON}

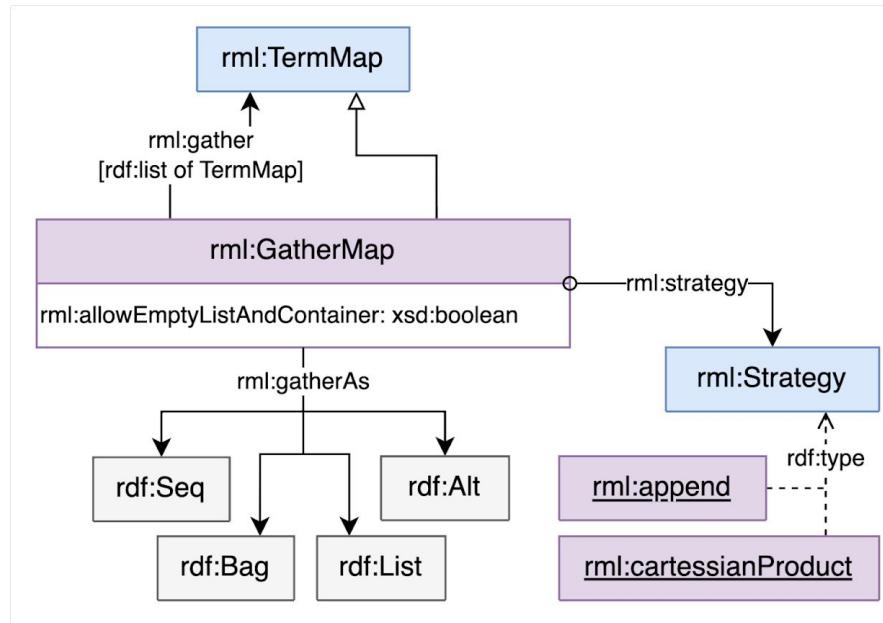
```
[ { "NAME": "Duplantis",
  "RANK": "1",
  "MARK": "6.22",
  "DATE": "02-25-2023"
}, {
  "NAME": "Guttormsen",
  "RANK": "2",
  "MARK": "6.00",
  "DATE": "03-10-2023} ]
```

RDF

```
:Duplantis a ex:Athlete ;
  ex:jumpsOnDate "2023-02-25".
:Guttormsen a ex:Athlete ;
  ex:jumpsOnDate "2023-10-03".
```

# RML-CC: Collections and Containers

---



- Introduces generation of **collections and containers**
- Specifies how **gather terms** into a CC and **manage** them: to assign them a IRI or BN, manage empty CC, how the gathering is performed...



[w3id.org/rml/cc](http://w3id.org/rml/cc)

[w3id.org/rml/cc/spec](http://w3id.org/rml/cc/spec)

[w3id.org/rml/cc/shapes](http://w3id.org/rml/cc/shapes)

# RML-CC: Example

---

---

```
<#RankingTriplesMap> a rml:TriplesMap ;
    rml:logicalSource <#JSONSource> ;
    rml:subjectMap [
        rml:constant :Ranking23 ];
    rml:predicateObjectMap [
        rml:predicate ex:contains;
        rml:objectMap [
            rml:gather ( [
                rml:template "{$.*.NAME}";
                rml:termType rml:IRI ] );
            rml:gatherAs rdf:List; ] ] .
```



{JSON}

```
[ { "NAME": "Duplantis",
  "RANK": "1",
  "MARK": "6.22",
  "DATE": "02-25-2023
}, {
  "NAME": "Guttormsen",
  "RANK": "2",
  "MARK": "6.00",
  "DATE": "03-10-2023} ]
```

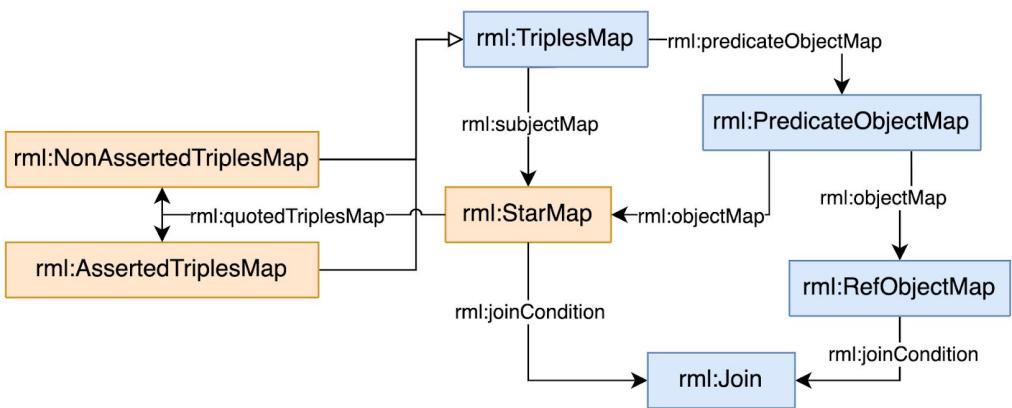


RDF

```
:Ranking23 ex:contains (:Duplantis
:Guttormsen ).
```

# RML-star: RDF-star

---



- **Recursiveness** in mapping rules to generate quoted triples
- Applicable in **subject** and **object** position
- Asserted and non-asserted quoted triples



[w3id.org/rml/star](http://w3id.org/rml/star)

[w3id.org/rml/star/spec](http://w3id.org/rml/star/spec)

[w3id.org/rml/star/shapes](http://w3id.org/rml/star/shapes)

# RML-star: Example

---

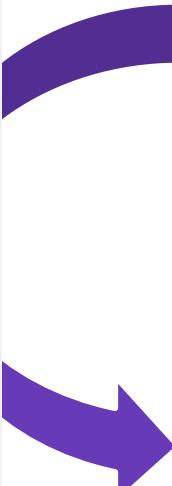
---

```
<#QuotedRankTM> a rml:AssertedTriplesMap ;
    rml:logicalSource <#JSONSource> ;
    rml:subjectMap [rml:template "{$NAME}"] ;
    rml:predicateObjectMap [
        rml:predicate ex:mark ;
        rml:objectMap [
            rml:reference "$.MARK" ;
        ] ] .

<#DateTM> a rml:TriplesMap ;
    rml:logicalSource <#JSONSource> ;
    rml:subjectMap [
        rml:quotedTriplesMap <#QuotedRankTM>
    ] ;
    rml:predicateObjectMap [
        rml:predicate ex:date ;
        rml:objectMap [
            rml:reference "$.DATE" ;
        ] ] .
```

{JSON}

```
[ { "NAME": "Duplantis",
    "RANK": "1",
    "MARK": "6.22",
    "DATE": "02-25-2023
  },
  {
    "NAME": "Guttormsen",
    "RANK": "2",
    "MARK": "6.00",
    "DATE": "03-10-2023 } ]
```



 :Duplantis ex:mark "6.22" .
<< :Duplantis ex:mark "6.22" >>
 ex:date "02-25-2023" .
:D guttormsen ex:mark "6.00" .
<< :Guttormsen ex:mark "6.00" >>
 ex:date "03-10-2023" .

# How do you construct a Knowledge Graph?

David Chaves-Fraga and  
Ana Iglesias-Molina

## Coffee Break!

We will be back at 16:00

# Easy KG construction

16:00 - 17:00

YARRRML: User-Friendly  
Syntax for RML

Hands-on example



— — —

# User-friendly approaches to write mappings

---

---

- [R2]RML mappings are written as Turtle files:
  - Risk of syntax errors
  - Too many language constructs
  - Long process (~6 months)
- Approaches developed to aid the mapping writing:
  - User-friendly syntax:
    - YARRRML → RML
    - XRM → R2RML, CSVW, RML (Zazuko)
    - SMS2 → R2RML (Stardog)
    - obda → R2RML (Ontop)
  - Editors
    - RMLEditor → RML
    - Mapeathor → R2RML, RML, YARRRML
    - ...

# YARRRML

---



User-friendly serialization for RML based on YAML syntax:

- Compact syntax
- Widely adopted in real-life projects
  - Entity Reconciliation API on Google Specification
  - <https://w3id.org/kg-construct/yarrrml>
- Online service
  - <https://rml.io/yarrrml/matey/>

# Mapping example

—

## YARRRML

## Data sources

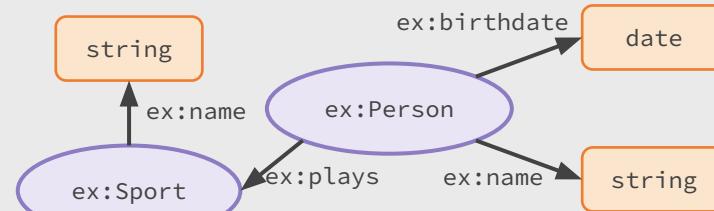
people.csv

ID	Name	Birthdate	SportID
1	Emily	08/02/90	2
2	Jonah	18/11/75	2

sports.csv

ID	Sport
1	Ice Skating
2	Rugby

## Ontology



## Mapping



Group of  
transformation rules

mappings:  
**PERSON:**

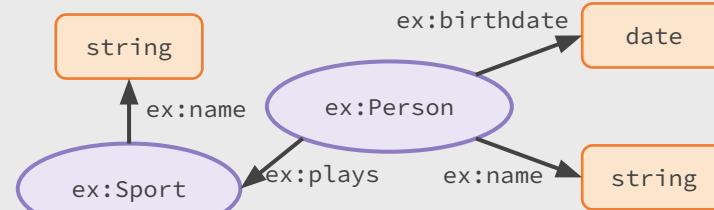
## Output Knowledge Graph

## Data sources

people.csv			
ID	Name	Birthdate	SportID
1	Emily	08/02/90	2
2	Jonah	18/11/75	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



## Mapping



```
mappings:  
PERSON:  
sources:  
- ['data/people.csv~csv']
```

## Output Knowledge Graph

## Data sources

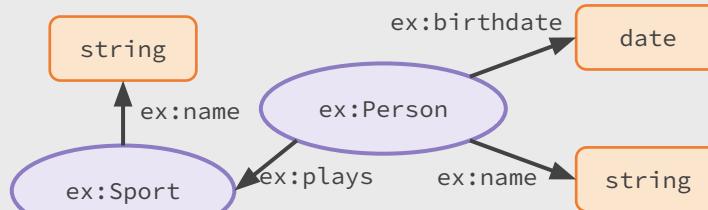
people.csv

ID	Name	Birthdate	SportID
1	Emily	08/02/90	2
2	Jonah	18/11/75	2

sports.csv

ID	Sport
1	Ice Skating
2	Rugby

## Ontology



## Mapping



```
mappings:  
PERSON:  
sources:  
- ['data/people.csv~csv']  
s: http://ex.com/Person/\$\(ID\)
```

## Output Knowledge Graph

<Person/1>

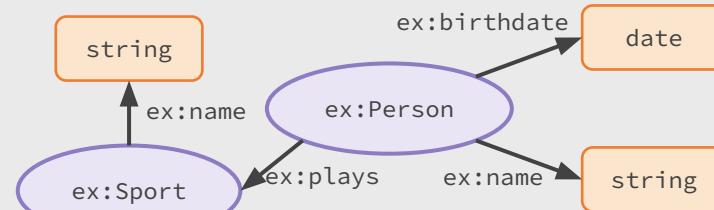
<Person/2>

## Data sources

people.csv			
ID	Name	Birthdate	SportID
1	Emily	08/02/90	2
2	Jonah	18/11/75	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology

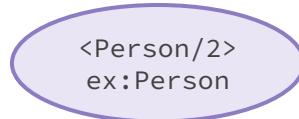


## Mapping



```
mappings:  
PERSON:  
sources:  
- ['data/people.csv~csv']  
s: http://ex.com/Person/${ID}  
po:  
- [a, ex:Person]
```

## Output Knowledge Graph

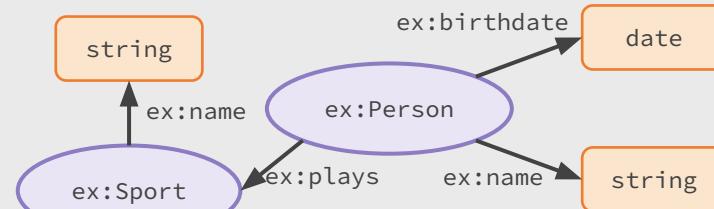


## Data sources

people.csv			
ID	Name	Birthdate	SportID
1	Emily	08/02/90	2
2	Jonah	18/11/75	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



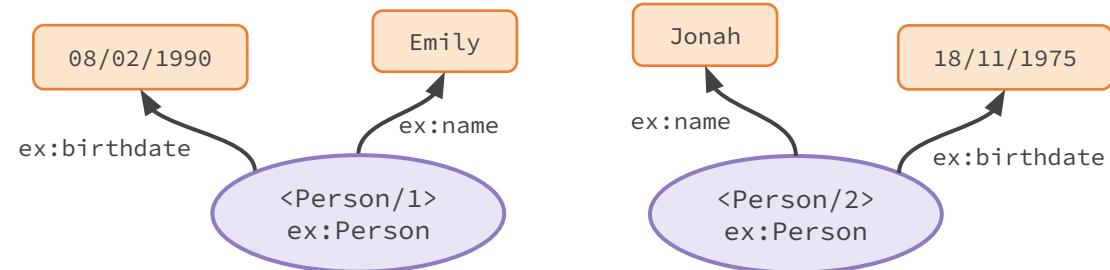
## Mapping



```

mappings:
PERSON:
sources:
- ['data/people.csv~csv']
s: http://ex.com/Person/${ID}
po:
- [a, ex:Person]
- [ex:name, ${Name}]
- [ex:birthdate, ${Birthdate}]
  
```

## Output Knowledge Graph

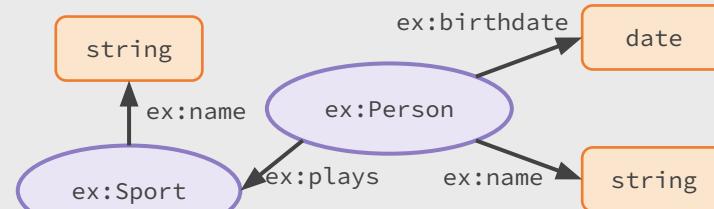


## Data sources

people.csv			
ID	Name	Birthdate	SportID
1	Emily	08/02/90	2
2	Jonah	18/11/75	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



## Mapping



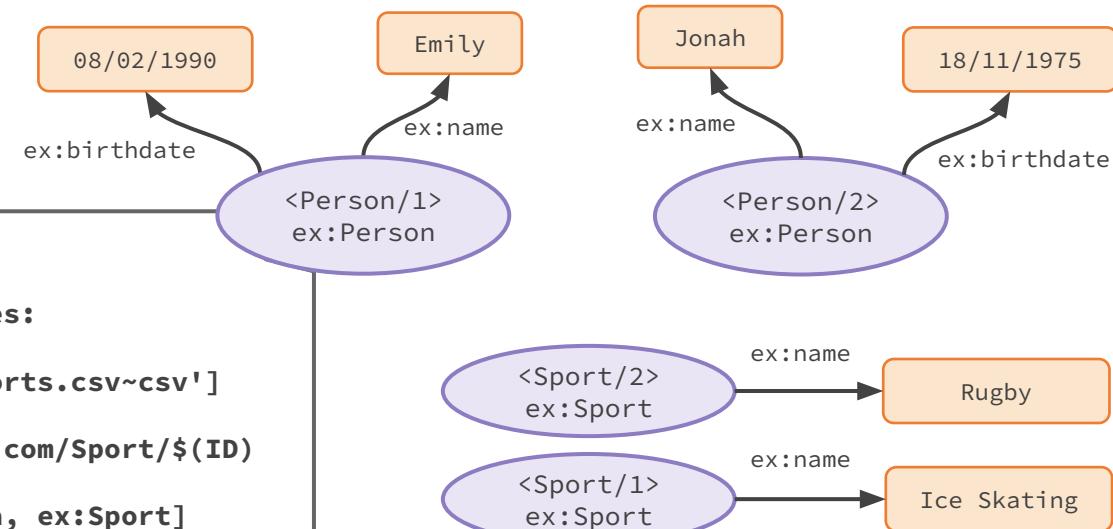
```

mappings:
PERSON:
sources:
- ['data/people.csv~csv']
s: http://ex.com/Person/${ID}
po:
- [a, ex:Person]
- [ex:name, ${Name}]
- [ex:birthdate, ${Birthdate}]
  
```

```

mappings:
SPORT:
sources:
- ['data/sports.csv~csv']
s:
http://ex.com/Sport/${ID}
po:
- [a, ex:Sport]
- [ex:name, ${Sport}]
  
```

## Output Knowledge Graph

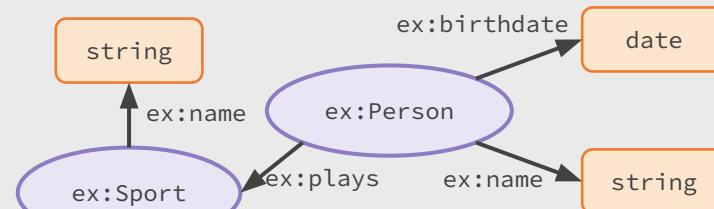


## Data sources

people.csv			
ID	Name	Birthdate	SportID
1	Emily	08/02/90	2
2	Jonah	18/11/75	2

sports.csv	
ID	Sport
1	Ice Skating
2	Rugby

## Ontology



## Mapping



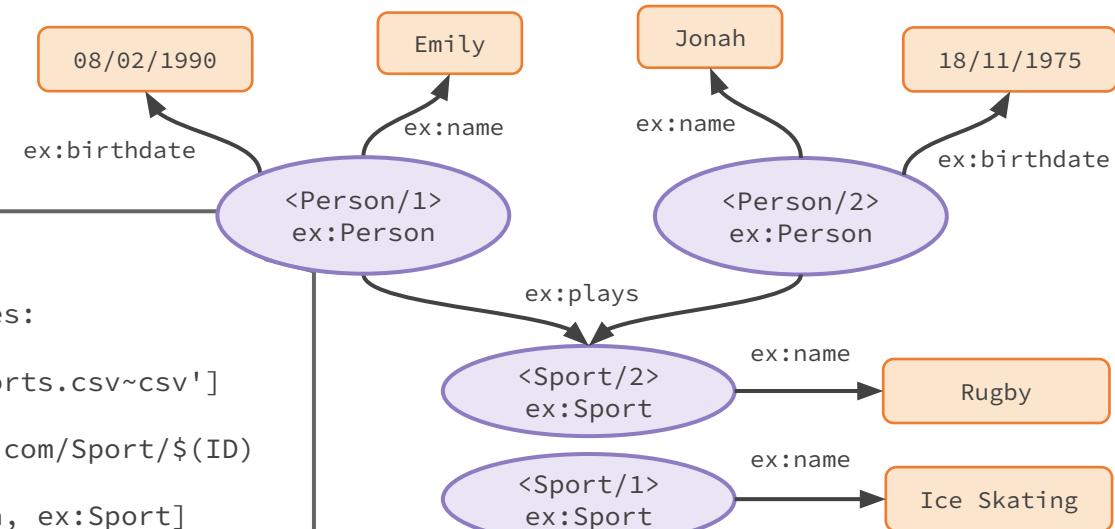
```

mappings:
PERSON:
sources:
- ['data/people.csv~csv']
s: http://ex.com/Person/${ID}
po:
- [a, ex:Person]
- [ex:name, ${Name}]
- [ex:birthdate, ${Birthdate}]
- p: ex:plays
o:
- mapping: SPORT
  condition:
    function: equal
  parameters:
    - [str1, ${SportID}]
    - [str2, ${ID}]
  
```

mappings:  
SPORT:

sources:  
- ['data/sports.csv~csv']  
s:  
http://ex.com/Sport/\${ID}  
po:  
- [a, ex:Sport]  
- [ex:name, \${Sport}]

## Output Knowledge Graph



# Hands-on exercise

# Hands-on exercise

---

- Creating a KG from a CSV file using YARRMML



# Hands-on exercise

---

- Creating a KG from a CSV file using YARRRML:
  1. Access the resources in: <https://bit.ly/4h2rgcW>
  2. Create the mapping completing the file  
‘yarrmrl\_template.yml’
  3. Can be tested using [Matey](#) or [Yatter](#)
  4. Run the mapping and the data with Morph-KGC to  
create the KG

# Mapping Supporters: Matey (YARRRML-parser)

The screenshot shows a web-based tool for generating RDF Mappings (RML) from YARRRML configurations. The interface includes tabs for People (JSON), Advanced, Facebook, Targets, Tutorial, Basic LDES, Actions, and various output formats like Generate RML, Generate LD, Reset state, and Layout.

**Input: Data**

```
1 palabra
2 comadreja
3 incendio
```

**Input: YARRRML**

```
1 prefixes:
2   emofind: "https://www.usc.gal/pcc/app/emofinder/"
3   emoskos: "https://www.usc.gal/pcc/app/emofinder/kos/"
4   idlab-fn: "http://example.com/idlab/function/"
5   grel: "http://users.ugent.be/~bjdmeest/function/grel.ttl#"
6
7 mappings:
8   palabras:
9     sources:
10      - ['data.csv~csv']
11      s: emofind:recurso/word/${palabra}
12      po:
13        - [a, emofind:Word]
14        - [emofind:name, ${palabra}, es-lang]
15      bases:
16        sources:
```

**Output: RML**

```
1 @prefix rr: <http://www.w3.org/ns/r2rml#> .
2 @prefix rml: <http://semweb.mmlab.be/ns/rml#> .
3 @prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
4 @prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#> .
5 @prefix ql: <http://semweb.mmlab.be/ns/ql#> .
6 @prefix map: <http://mapping.example.com/> .
7 @prefix as: <https://www.w3.org/ns/activitystreams#> .
8 @prefix skos: <http://www.w3.org/2004/02/skos/core#> .
9 @prefix xsd: <http://www.w3.org/2001/XMLSchema#> .
10 @prefix sd: <http://www.w3.org/ns/sparql-service-description#> .
11 @prefix emofind: <https://www.usc.gal/pcc/app/emofinder/> .
12 @prefix emoskos: <https://www.usc.gal/pcc/app/emofinder/kos/> .
13 @prefix idlab-fn: <http://example.com/idlab/function/> .
14 @prefix grel: <http://users.ugent.be/~bjdmeest/function/grel.ttl#> .
15
16 map:fn_000 rml:logicalSource map:source_002 :
```

**Output: Turtle/TriG**

```
1 @prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
2 @prefix os: <https://www.w3.org/ns/activitystreams#> .
3 @prefix skos: <http://www.w3.org/2004/02/skos/core#> .
4 @prefix xsd: <http://www.w3.org/2001/XMLSchema#> .
5 @prefix sd: <http://www.w3.org/ns/sparql-service-description#> .
6 @prefix emofind: <https://www.usc.gal/pcc/app/emofinder/> .
7 @prefix emoskos: <https://www.usc.gal/pcc/app/emofinder/kos/> .
8 @prefix idlab-fn: <http://example.com/idlab/function/> .
9 @prefix grel: <http://users.ugent.be/~bjdmeest/function/grel.ttl#> .
10
11 <https://www.usc.gal/pcc/app/emofinder/recurso/word/comadreja> rdf:type emofind:Word .
12   emofind:name "comadreja"@es .
13
14 <https://www.usc.gal/pcc/app/emofinder/recurso/word/incendio> rdf:type emofind:Word .
15   emofind:name "incendio"@es .
16
```

## Contact

yarrmml@rml.io

<https://rml.io/yarrmml/matey/> and <https://github.com/RMLio/yarrmml-parser> (NodeJS)

# Hands-on exercise

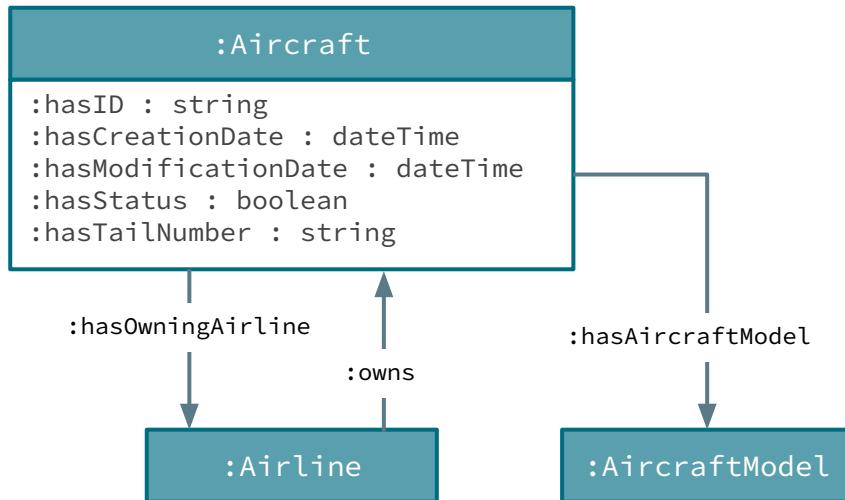
---

---

- Creating a KG from a CSV file using YARRMML



<https://bit.ly/4h2rgcW>



**sfo-aircraft-tail-numbers-and-models.csv**

- aircraft\_id
- aircraft\_model
- airline
- creation\_date
- modification\_date
- status
- tail\_number

# Resources for KG Construction

17:00 - 17:25

Ontology development

Mapping supporters

KG Construction engines

Benchmarks





# The RML ecosystem

## Compliant systems

- [RMLMapper](#)
- [RocketRML](#)
- [Chimera](#)
- [Morph-KGC](#)
- [SDM-RDFizer](#)
- [CARML](#)
- [RDF Processing Toolkit](#)
- [Mapeathor](#)
- [FlexRML](#)
- [BURP](#)
- [MEL](#)
- [Excel in RML](#)
- [Matey](#)
- [Yatter](#)
- [Spread2RML](#)
- [Dragoman](#)
- [RMLEditor](#)
- [YARRRML-parser](#)
- [mapping-template](#)
- [RMLStreamer](#)
- [RMLWeaver-js](#)

## Interoperability

- [SPARQL-Generate](#)
- [ShExML](#)
- [Helio](#)
- [YARRRML](#)
- [XRM](#)

## Benchmarks

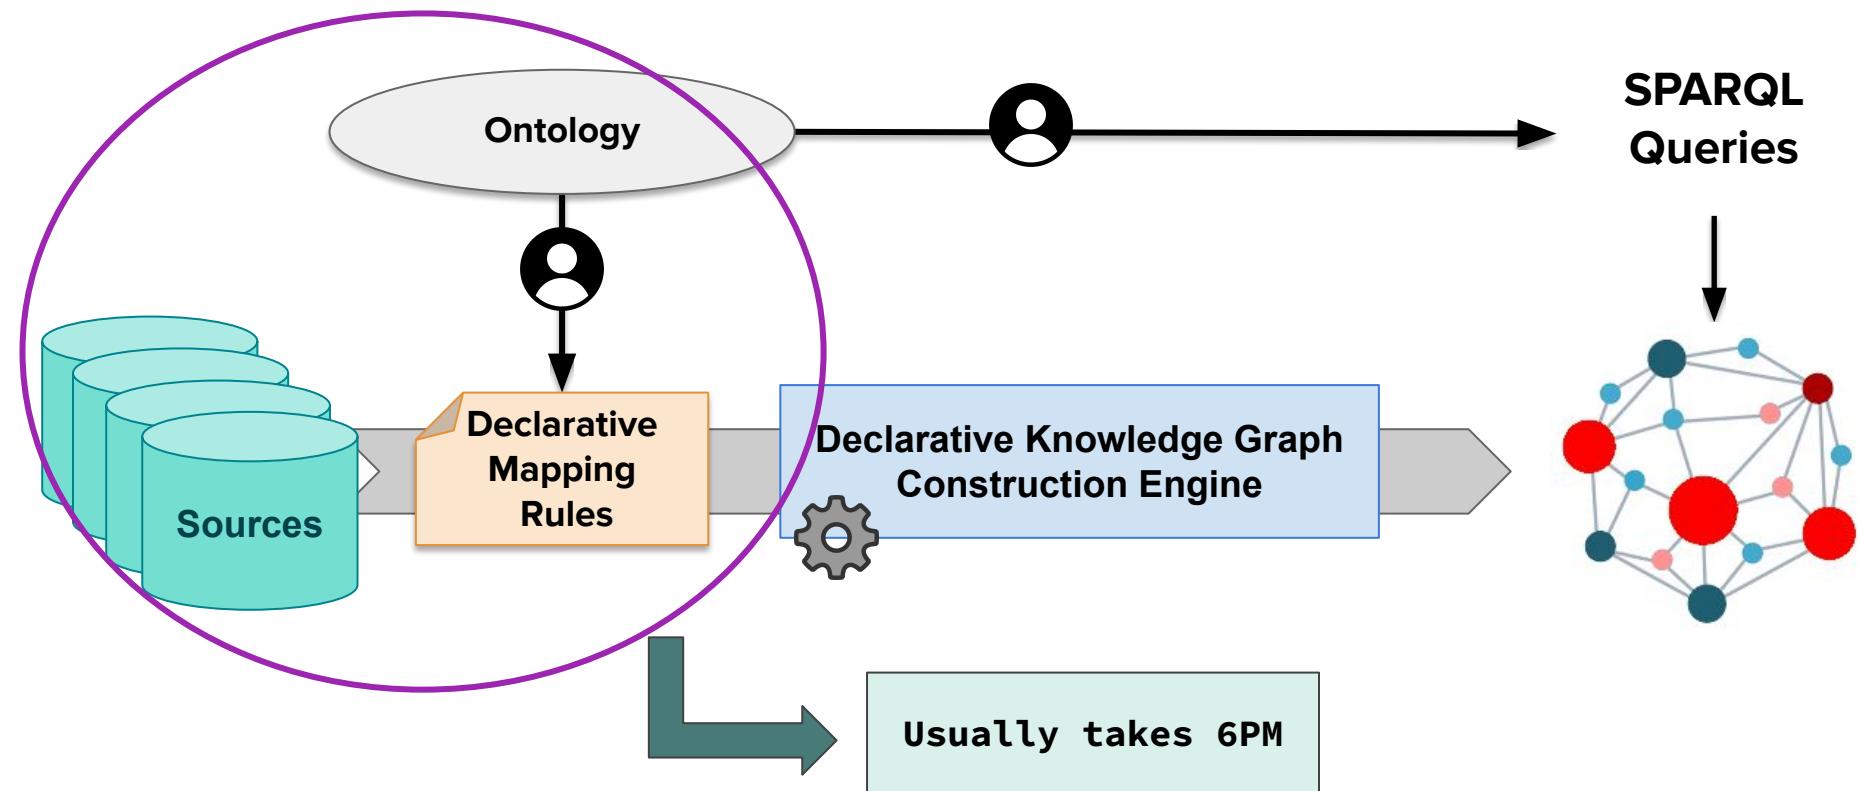
- [GTFS-Madrid Bench](#)
- [SDM-Genomic benchmark](#)
- [LUBM4OBDA](#)
- [KROWN \(KGC Challenge\)](#)

## Use cases

- National and European Projects
- European Commission (ERA, PPDS)
- Companies (Orange, BASF, REALE, Bosch, Google)

# Mapping rules

---



# Mapping Supporters: Yatter

---

---

## YATTER

license Apache-2.0 DOI 10.5281/zenodo.7024501 pypi v1.1.5 release date august codecov 83%

The tool translates mapping rules from YARRRML in a turtle-based serialization of RML or R2RML.

### Installation:

```
pip install yatter
```



### Execution from CLI

To execute from command line run the following:

- From YARRRML to [R2]RML

```
python3 -m yatter -i path_to_input_yarrmml.yml -o path_to_rdf_mapping.ttl [-f R2RML]
```



- From [R2]RML to YARRRML

```
python3 -m yatter -i path_to_input_rdf_mapping.ttl -o path_to_output_yarrmml.yml [-f R2RML]
```



`-f R2RML` is an optional parameter for translating input YARRRML to R2RML (and inverse)

- Easy to read translation
- Python-based
- Code:

<https://github.com/oeg-upm/yatter>



Iglesias-Molina, A., Chaves-Fraga, D., Dasoulas, I., & Dimou, A. (2023). Human-friendly RDF graph construction: Which one do you chose?. In International Conference on Web Engineering



# Mapping Supporters: Mapeathor

Prefix sheet		Subject sheet			
Prefix	URI	ID	Class	URI	Graph
ex	http://ex.com/	PERSON	ex:Person	http://ex.com/Person/{ID}	http://ex.com/graph
fno	https://w3id.org/function/ontology#	SPORT	ex:Sport	http://ex.com/Sport/{ID}	
grel	http://semweb.dasciencelab.be/ns/grel#	SPORT	ex:TeamSport	http://ex.com/Sport/{ID}	

Source sheet			Function sheet		
ID	Feature	Value	FunctionID	Feature	Value
PERSON	source	data/people.csv	<Fun1>	executes	grel:toDate
PERSON	format	CSV	<Fun2>	returns	grel:output_date
SPORT	source	data/sports.csv	<Fun3>	exec:valueParam1	yyyyMMdd
SPORT	format	CSV	<Fun4>	exec:valueParam2	d/M/y

Predicate_Object sheet							
ID	Predicate	Object	DataType	Language	ReferenceID	InnerRef	OuterRef
PERSON	ex:name	{name}	string	en			
PERSON	ex:birthdate	{birthdate}	date				
PERSON	ex:sport				SPORT	{SportID}	{ID}
SPORT	ex:name	{sport}	string	en			
SPORT	ex:code	{ID}	integer				
SPORT	ex:comment	<Fun5>					

## Mapeathor Demo

You can test this tool with the demo below, just use a Google Spreadsheet URL or an Excel file (XLSX) and select a mapping format. Check the templates and examples in the GitHub repository. You can also check [this Swagger instance](#).

Choose the output format:

RML

R2RML

YARRRML

Do you prefer upload a XLSX file?

Google Spreadsheet URL:

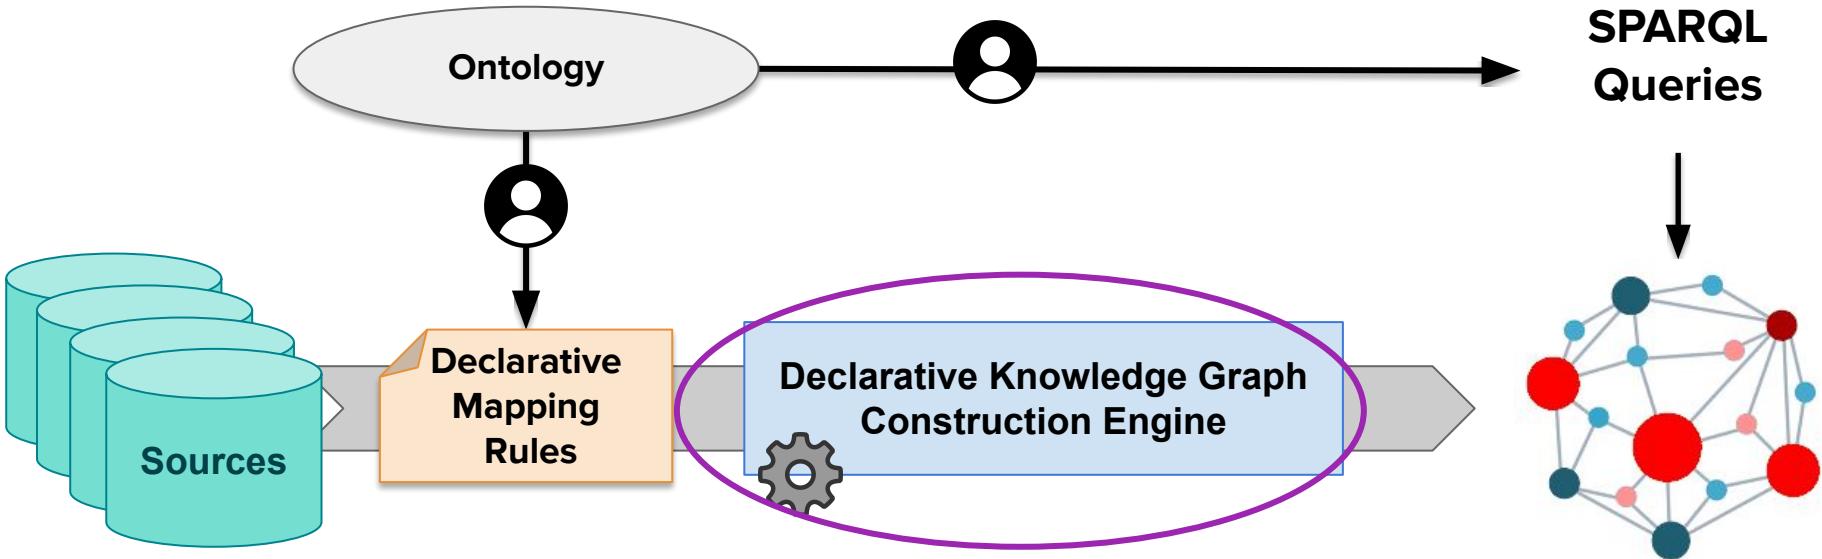
- Excel-based mapping
- Translation into YARRRML, R2RML and RML
- REST API, Python Library
- Code: <https://github.com/oeg-upm/mapeathor>



Iglesias-Molina, A., Pozo-Gilo, L., Dona, D., Ruckhaus, E., Chaves-Fraga, D., & Corcho, O. (2020). Mapeathor: Simplifying the Specification of Declarative Rules for Knowledge Graph Construction. In International Semantic Web Conference

# Mapping rules

---



# Engines for KG Construction: SDM-RDFizer

---

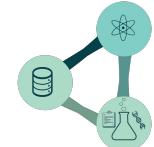
---

Install at: <https://pypi.org/project/rdfizer/> (pip install rdfizer)

Github website: <https://github.com/SDM-TIB/SDM-RDFizer>

## Features:

- Very efficient for large datasets
- Efficient execution of joins and duplicates removal
- Python based
- A bit difficult to run (config file needs to be filled)
- Some quality problems



SDM-RDFizer

# Engines for KG Construction: RMLMapper

---

Download/Install at: <https://github.com/RMLio/rmlmapper-java>

## Features:

- Good parser in terms on data quality
- Java based
- Really simple to run (-d is for removing duplicates):  
**java -jar rmlmapper.jar -m mapping.rml -o output.nt -d**
- Not very efficient when you have many joins or duplicates



# Engines for KG Construction: RMLStreamer

---

Download/Install at: <https://github.com/RMLio/RMLStreamer>

## Features:

- To mapping stream data into RDF streams
- Scala based
- Not very efficient when you have joins



# Engines for KG Construction: FlexRML

---

KGC Challenge 2024

T2: Performance Award



THIS CERTIFICATE IS PROUDLY PRESENTED TO:

**FlexRML**

*FlexRML* is the fastest and most resource-efficient engine across Track 2  
of 2<sup>nd</sup> edition the Knowledge Graph Construction Workshop Challenge

Download/Install at: <https://github.com/wintechis/flex-rml>

Install: <https://www.npmjs.com/package/flexrml-node>

## Features:

- Only supports CSV data
- Super efficient (the best at this moment)
- It can runs materialization on the edge
- Written in C++ (JavaScript version also available)

# Engines for KG Construction: Morph-KGC

---

Install at: <https://pypi.org/project/morph-kgc/> (pip install morph-kgc)

Documentation: <https://morph-kgc.readthedocs.io/en/latest/documentation/>

## Features:

- Very efficient for big data
- Python based
- Optimizations in planification of the mappings
- Continuous integration for testing
- Generates RDF-star datasets

The logo consists of the word "morp" in a lowercase sans-serif font. The letters are colored blue, red, and blue again, with the 'o' being red.

# Let's play with Yatter+Morph-KGC in our example

---

## > [Yatter Tutorial](#)

This tool translates mapping rules from YARRRML in a turtle-based serialization of RML or R2RML.

[ ] ↳ 6 celdas ocultas

## > [Morph-KGC Tutorial](#)

[Morph-KGC](#) is an engine that constructs [RDF](#) and [RDF-star](#) knowledge graphs from heterogeneous data sources with the [R2RML](#) and [RML](#) mapping languages. The full documentation of Morph-KGC is available in [Read the Docs](#).

There are two different options to run Morph-KGC:

- As a [library](#), integrating with [RDFLib](#) and [Oxigraph](#).
- Via the [command line](#).

This tutorial shows the different alternatives to run Morph-KGC. Here, we use [RML](#) mappings, but the more user-friendly [YARRRML](#) mapping format is also supported. Data transformation, computation, or filtering before generating triples is also supported with [RML-FNML](#).

▶ ↳ 24 celdas ocultas

# Benchmarks for KG Construction

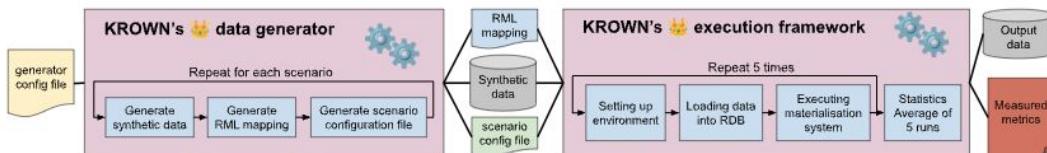
---

## KROWN 🔥: A Benchmark for RDF Graph Materialization

license [MIT](#) DOI [10.5281/zenodo.10979322](#) website [KROWN](#) docs [Data generator](#) docs [Execution framework](#)

KROWN 🔥 is a benchmark for materialization systems to construct Knowledge Graphs from (semi-)heterogeneous data sources using declarative mappings such as [RML](#).

Many benchmarks already exist for virtualization systems e.g. [GTFS-Madrid-Bench](#), [NPD](#), [BSBM](#) which focus on complex queries with a single declarative mapping. However, materialization systems are unaffected by complex queries since their input is the dataset and the mappings to generate a Knowledge Graph. Some specialized datasets exist to benchmark specific limitations of materialization systems such as duplicated or empty values in datasets e.g. [GENOMICS](#), but they do not cover all aspects of materialization systems. Therefore, it is hard to compare materialization systems among each other in general which is where KROWN 🔥 comes in!



<https://github.com/kg-construct/krown>



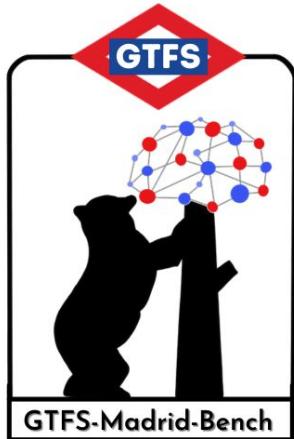
Van Assche, D., Chaves-Fraga, D., & Dimou, A. KROWN: A Benchmark for RDF Graph Materialisation. International Semantic Web Conference (2024)

# Benchmarks for KG Construction

---

## The GTFS-Madrid-Bench

We present GTFS-Madrid-Bench, a benchmark to evaluate declarative KG construction engines that can be used for the provision of access mechanisms to (virtual) knowledge graphs. Our proposal introduces several scenarios that aim at measuring performance and scalability as well as the query capabilities of all this kind of engines, considering their heterogeneity. The data sources used in our benchmark are derived from the [GTFS](#) data files of the subway network of Madrid. They can be transformed into several formats (CSV, JSON, SQL and XML) and scaled up. The query set aims at addressing a representative number of SPARQL 1.1 features while covering usual queries that data consumers may be interested in.



<https://github.com/oeg-upm/gtfs-bench>



Chaves-Fraga, D., Priyatna, F., Cimmino, A., Toledo, J., Ruckhaus, E., & Corcho, O. (2020). GTFS-Madrid-Bench: A benchmark for virtual knowledge graph access in the transport domain. *Journal of Web Semantics*.

# Benchmarks for KG Construction

---

SDM-GENOMICS Dataset: <https://doi.org/10.57702/4c9ivpgs>

- Testbeds for duplicate removal and join conditions
- RML support

BSBM: <http://wbsg.informatik.uni-mannheim.de/bizer/berlinsparqlbenchmark/>

- Well-known benchmark to test triplestores capabilities
- Adapted to virtual KG Construction with R2RML

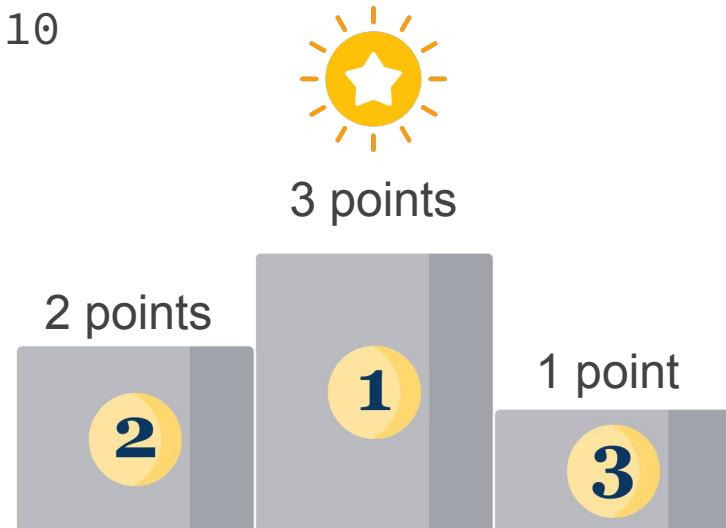
NPD: <https://github.com/ontop/npd-benchmark>

- Focuses on reasoning during virtual KG Construction
- R2RML suport

# Last performance results: GTFS-Madrid-Bench + KG parameters

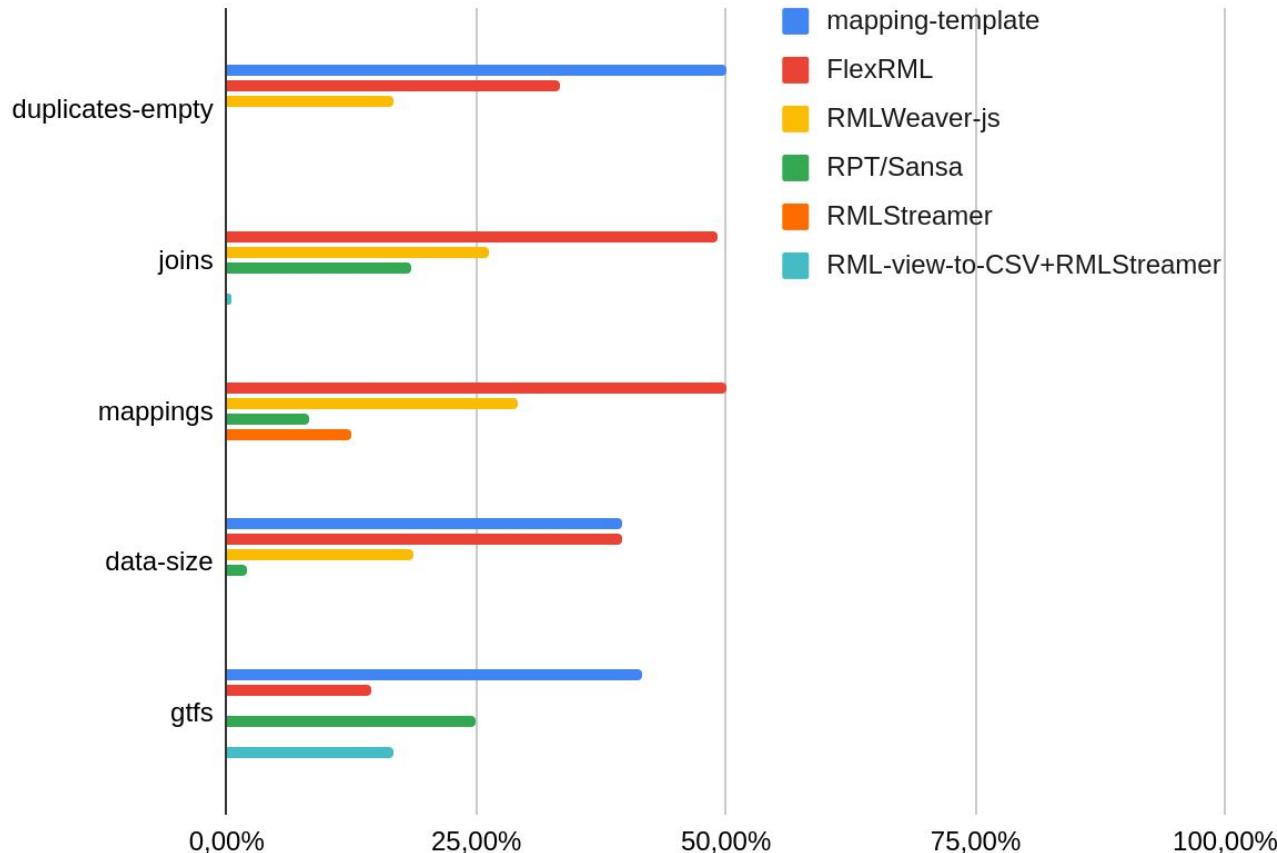
---

- **3 Variables:** execution time, CPU and RAM usage
- **65 Scenarios:**
  - KG parameters: 57
    - Duplicates & empty values: 10
    - Data size: 8
    - Mappings: 4
    - Joins: 35
  - GTFS: 8
- **Who gets the award?**
  - Top 3 tools per scenario and variable get points
  - Final count of all points



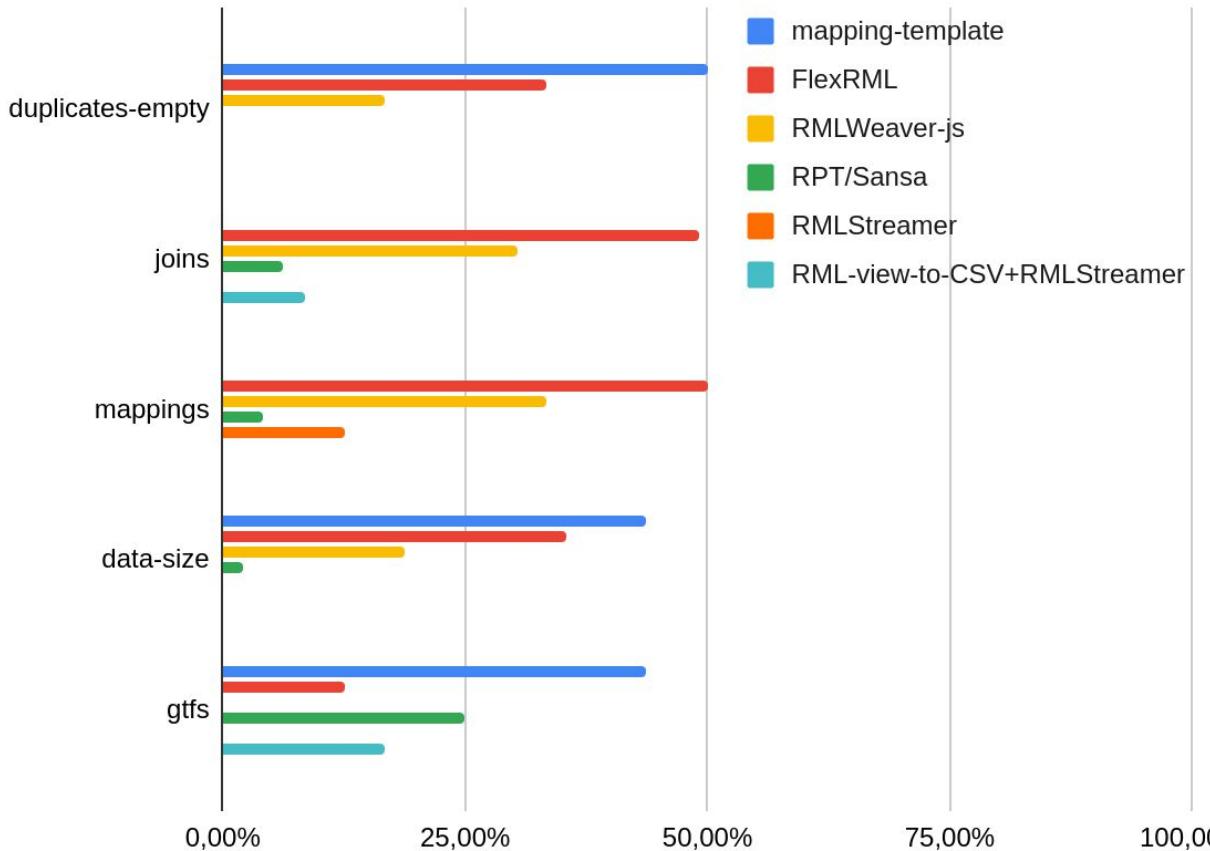
# Track 2: Points for execution time

---



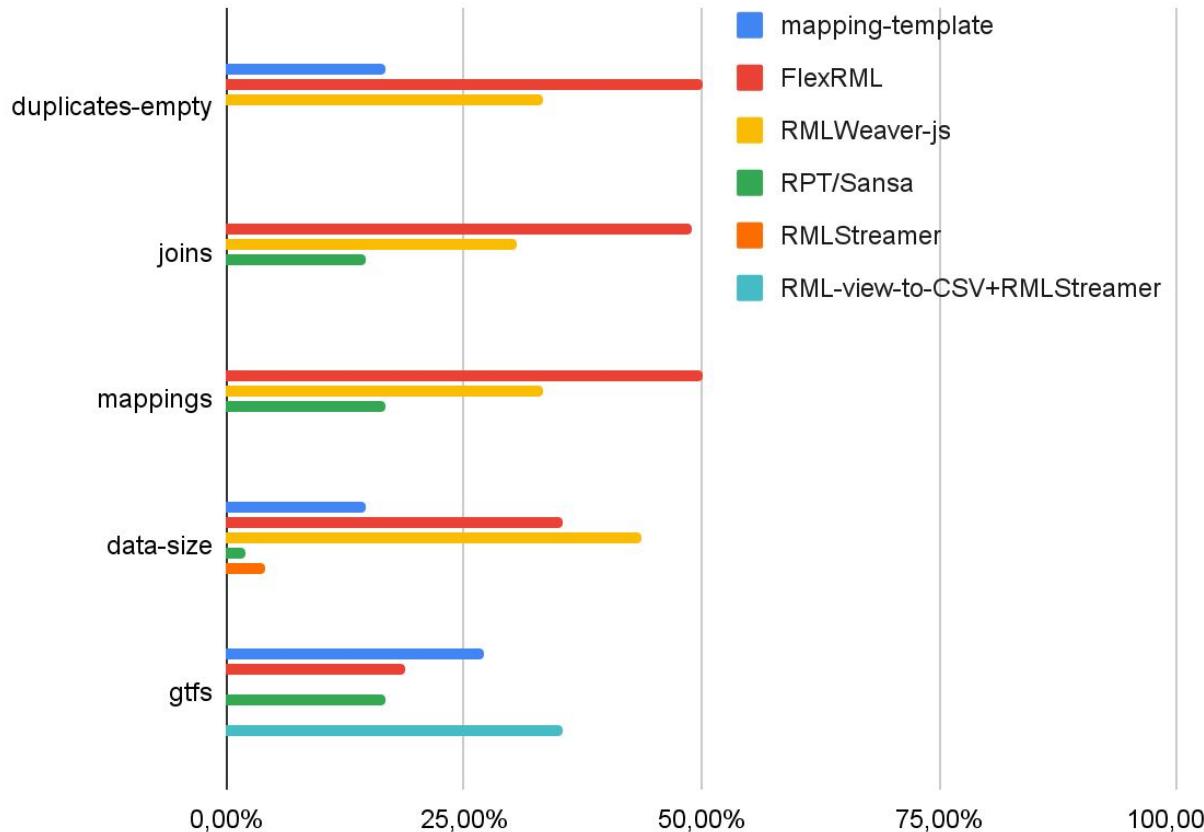
# Track 2: Points for CPU usage

---



# Track 2: Points for RAM usage

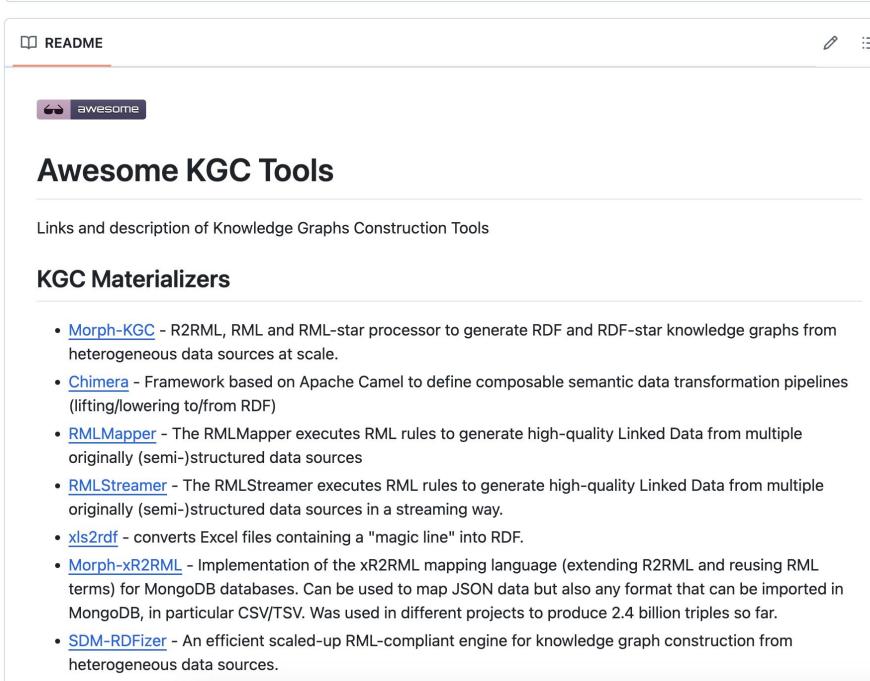
---



# Other resources: AWESOME-KGC-TOOLS list

---

---



The screenshot shows a GitHub README page for the "awesome-kgc-tools" repository. The page has a dark theme with a purple header bar. At the top left is a "README" button. On the right are edit and more options icons. Below the header is a purple "awesome" badge with a brain icon. The main title is "Awesome KGC Tools". A subtitle below it reads "Links and description of Knowledge Graphs Construction Tools". A section titled "KGC Materializers" contains a bulleted list of tools:

- [Morph-KGC](#) - R2RML, RML and RML-star processor to generate RDF and RDF-star knowledge graphs from heterogeneous data sources at scale.
- [Chimera](#) - Framework based on Apache Camel to define composable semantic data transformation pipelines (lifting/lowering to/from RDF)
- [RMLMapper](#) - The RMLMapper executes RML rules to generate high-quality Linked Data from multiple originally (semi-)structured data sources
- [RMLStreamer](#) - The RMLStreamer executes RML rules to generate high-quality Linked Data from multiple originally (semi-)structured data sources in a streaming way.
- [xls2rdf](#) - converts Excel files containing a "magic line" into RDF.
- [Morph-xR2RML](#) - Implementation of the xR2RML mapping language (extending R2RML and reusing RML terms) for MongoDB databases. Can be used to map JSON data but also any format that can be imported in MongoDB, in particular CSV/TSV. Was used in different projects to produce 2.4 billion triples so far.
- [SDM-RDFizer](#) - An efficient scaled-up RML-compliant engine for knowledge graph construction from heterogeneous data sources.

<https://kg-construct.github.io/awesome-kgc-tools/>

# Ontology Development

---

---

Sunday 20th (tomorrow)

09:00 - 12:30

Room: TBC

<https://chowlk.linkeddata.es/>

<https://lot.linkeddata.es/>

## Boosting ontology conceptualization with Chowlk

Tutorial at [ECAI 2024](#)

### 1. Description

This tutorial aims to support ontology developers to reduce the time consumed in ontology implementation by generating graphical ontology conceptualizations that can be converted into code automatically. Participants will gain knowledge about the Chowlk visual notation for ontologies, how to use the Chowlk converter to accelerate ontology implementation based on the generated conceptualizations, how to use the related resources as templates and, finally, how to create new templates.

This tutorial will focus on explaining 1) the Chowlk notation, origin, motivation and current status; 2) how to use the converter to accelerate the ontology implementation activity based on the generated diagram; 3) how to use the related resources as templates; and finally 4) how to create new templates. The tutorial will be mainly practical and interactive, driving the attendees through the Chowlk framework features by using them. In this sense we will "Introduce expert non-specialists to an AI subarea" (ontology development) and "Provide instruction in established but specialized AI methodologies."

### 2. Speakers

María Poveda-Villalón

[mpoveda \(at\) fi.upm.es](mailto:mpoveda(at)fi.upm.es)

Dr. María Poveda-Villalón is an associate professor at the Artificial Intelligence Department of the Universidad Politécnica de Madrid and is also part of the Ontology Engineering Group research lab. Her research activities focus on Ontological Engineering, Ontology Evaluation, Knowledge Representation and the Semantic Web. Previously she finished her studies in Computer Science (2009) by Universidad Politécnica de Madrid, and then she moved to study the Artificial Intelligence Research Master finished in 2010 in the same university. She has contributed to the ontology engineering field developing methodologies and tools like OOPS! (Ontology Pitfall Scanner!) which has been broadly adopted by the community, the LOT methodology for building ontologies and the Chowlk framework for ontology conceptualization and implementation. She has worked in the context of Spanish research projects as well as in European. She has contributed to the organization of the "Linked Data in Architecture and Construction Workshop" since 2015 edition, the "13th OWL: Experiences and Directions Workshop and 5th OWL reasoner evaluation workshop" in 2016, the Workshop on Ontology Patterns (WOP) in 2019 and 2022, the "Linked Energy Data Vocamp" in 2015 and she has organized the 2nd Summer School of Linked Data in Architecture and Construction. Finally, she is part of the W3C Web of Things Working Group. Her teaching activities involve teaching semantic web and ontological engineering in different courses at official degrees, masters courses, MOOCs and online masters as well to deliver courses to the administration and private companies. She has carried out similar tutorials as "Catching up with ontological engineering. To git-commit and beyond" at EKAW2018 and "Integrating ontological development with software engineering trends" FDL2019.



# Conclusions

---

- Systematic construction of RDF Knowledge Graphs
- RML as declarative language for KGC
- YARRRML as human-syntax to facilitate the use
- Several engines with their own features
- Benchmarks to test performance (time, memory, CPU)
- Big community (>180 members in W3C)
- Towards the standardization process
- Impact over companies (Google, Orange, Bosch, etc)
- In-production on European Data Spaces (PPDS, ERA)

# How do you construct a Knowledge Graph?

David Chaves-Fraga and  
Ana Iglesias-Molina

Thanks for  
attending!